Learn from Home | Flat 50% OFF SITEWIDE | Use Coupon - WHIZSITE50

Practice Test III

Completed on 21-July-2020

Congratulations, you passed. Keep it up

Attempt	Marks Obtained	Your score	Time Taken	Result
02	49 / 55	89.09%	00 H 09 M 32 S	Congratulations! Passed

Share your Result with your friends

Domains wise Quiz Performance Report

No	Domain	Total Question	Correct	Incorrect	Unattempted	Marked as Review
1	Connect to and consume Azure services and third-party services	14	12	2	0	0
2	Monitor, troubleshoot, and optimize Azure solutions	9	8	1	0	0
3	Develop Azure compute solutions	14	12	2	0	0

4	Develop for Azure storage	18	17	1	0	0
Total	All Domain	55	49	6	0	0

Review the Answers

Sorting by All

Question 1

Correct

Domain :Connect to and consume Azure services and third-party services

A software company is developing a software solution. The software solution is for a food delivery-based company. The software needs to adhere to the following workflow

A driver selects the restaurants for which they will deliver orders.

Orders are sent to all available drivers in an area.

Only orders for the selected restaurants will appear for the driver.

The first driver to accept an order removes it from the list of available orders.

The application needs to make use of the Azure Service Bus service.
Which of the following actions would you implement for this requirement? Choose 3 answers from the options given below

A. **Create a Service Bus topic for each restaurant for which a driver can receive messages.**

✓ B. **Create a single Service Bus topic**

C. **Create a single Service Bus subscription**

✓ D. **Create a single Service Bus Namespace**

E. **Create a Service Bus Namespace for each restaurant for which a driver can receive messages.**

✓ F. **Create a Service Bus Subscription for each restaurant for which a driver can receive messages.**

Explanation:

Answer – B, D and F

You should first create a Service Bus Namespace. Option E is incorrect since creating a namespace for each restaurant would just be a maintenance overhead and difficult to keep track via a program.

Here since the driver needs to choose the restaurant, that means the driver can be a subscriber.

Here you should have just one Topic. If you have multiple topics, then an order needs to be sent to all topics. Then deleting an order once it has been picked by a driver will an issue. So, Option A gets ruled out.

You can create subscriptions and create rules based on driver and area.

For more information on Azure Service Bus, one can go to the below link

https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messaging-overview

Ask our Experts Rate this Question?

View Queries open

Question 2 Correct

Domain :Connect to and consume Azure services and third-party services

A company has an application that provides product data to external consultants.
Azure API Management is used to publish API's to the consultants.
The API needs to meet the following requirements

Support alternative input parameters.

Remove formatting text from responses.

Provide additional context to back-end services.

Which type of policy would you use for the following requirement
"Rewrite the request URL to match to the format expected by the web service"

A. **Inbound**

B. **Outbound**

C. **Backend**

D. **Error**

Explanation:

Answer – A

An example of this is given in the Microsoft documentation

Policy statement

XML Copy

```
<rewrite-uri template="uri template" copy-unmatched-params="true | false" />
```

Example

XML Copy

```
<policies>
    <inbound>
        <base />
        <rewrite-uri template="/v2/US/hardware/{storenumber}&{ordernumber}?City=city&State=state" />
    </inbound>
    <outbound>
        <base />
    </outbound>
</policies>
```

Since this is clearly mentioned, all other options are incorrect

For more information on API management transformation URL's, one can go to the below link

https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies

Ask our Experts Rate this Question?

View Queries open

Question 3 Correct

Domain :Connect to and consume Azure services and third-party services

A company has an application that provides product data to external consultants.
Azure API Management is used to publish API's to the consultants.
The API needs to meet the following requirements

Support alternative input parameters.

Remove formatting text from responses.

Provide additional context to back-end services.

Which type of policy would you use for the following requirement
"Remove formatting text from responses"

A. **Inbound**

✓ B. **Outbound**

C. **Backend**

D. **Error**

Explanation:

Answer - B

You can use policy expressions and find and replace to format the text in the response. An example of this is given in the Microsoft documentation

```
<policies>
  <inbound>
    <!-- Save the URL in invoked request from the context before its rewritten. -->
    <set-variable name="requestPath" value="@(context.Request.Url.Path)" />
    <base />
    <rewrite-uri template="api/add?a={a}&amp;b={b}" />
  </inbound>
  <outbound>
  <base />
    <!-- Insert the saved request method+URL in the result section. -->
    <!-- Note request method wasn't rewritten, so we can access that from context. -->
    <find-and-replace from="&lt;result&gt;" to="@{
      string str = "<result>\n    <operation>";
      str += context.Request.Method;
      str += " \"";
      str += context.Variables.GetValueOrDefault<string>("requestPath");
      str += "\"</operation>";
      return str;
    }" />
  </outbound>
</policies>
```

Since this is clearly mentioned, all other options are incorrect

For more information on API management transformation URL's and a blog article on the same, one can go to the below link

https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies

https://azure.microsoft.com/fr-fr/blog/policy-expressions-in-azure-api-management/

Ask our Experts Rate this Question?

View Queries open

Question 4 Correct

Domain :Connect to and consume Azure services and third-party services

A company has an application that provides product data to external consultants.
Azure API Management is used to publish API's to the consultants.
The API needs to meet the following requirements

Support alternative input parameters.

Remove formatting text from responses.

Provide additional context to back-end services.

Which type of policy would you use for the following requirement
"Forward the user ID that is associated with the subscription key for the original request to the back-end service"

A. **Inbound**

B. **Outbound**

C. **Backend**

D. **Error**

Explanation:

Answer - A

An example of this is given in the Microsoft documentation

XML Copy

```
<!-- The policies described in this file show how to send some context information to the backend service

<!-- Copy these snippets into the inbound element -->

<policies>
  <inbound>
    <base />
      <!-- Forward the name of the product associated with the subscription key in the request to the back
      <set-query-parameter name="x-product-name" exists-action="override">
        <value>@(context.Product.Name)</value>
      </set-query-parameter>

      <!-- Forward the user id associated with the subscription key in the request as well as the region w
      <set-header name="x-request-context-data" exists-action="override">
        <value>@(context.User.Id)</value>
        <value>@(context.Deployment.Region)</value>
      </set-header>
  </inbound>
  <backend>
```

Since this is clearly mentioned, all other options are incorrect

For more information on API management transformation URL's, one can go to the below link

https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies

Ask our Experts Rate this Question?

View Queries open

Question 5 Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A development team has published a Web App to the Azure Web App Service. They are also using Application Insights for the Web App for monitoring purposes. They have to ensure that the cost for Application Insights does not exceed a pre-set budget. Which of the following would you implement to adhere to this requirement?

A. **Implement ingestion sampling using the Azure portal.**

B. **Set a daily cap for the Application Insights instance.**

C. **Implement adaptive sampling using the Azure portal.**

D. **Implement adaptive sampling using the Application Insights SDK.**

E. **Implement ingestion sampling using the Application Insights SDK.**

Explanation:

Answer – D

Adaptive sampling is the default for the ASP.NET SDK. Adaptive sampling automatically adjusts to the volume of telemetry that your app sends. It operates automatically in the SDK in your web app so that telemetry traffic on the network is reduced.

Option B: Daily Cap is also possible answer, but it has a major disadvantage. In case data capacity is exhausted any point during day, then rest of the day, no data will be ingested and hence in case of any problem in the VM will not be notified. Hence this is not recommended.

For more information on managing costs for Application Insights, one can go to the below link

https://docs.microsoft.com/en-us/azure/azure-monitor/app/pricing#managing-your-data-volume

Ask our Experts Rate this Question?

View Queries open

Question 6 Incorrect

Domain :Monitor, troubleshoot, and optimize Azure solutions

A development team is developing an application. The application will be working with customer data. The application will also be making use of Azure Redis Cache. You need to invalidate the cache when the customer data is changed.
You have to complete the below code to comply with the requirement

```
void clearCustomerCache(string p_Customer)
{
        //Establish the cache connection
                Slot1
        //Invalidate the cache
                Slot2
}
```

Which of the following will go into Slot1?

A. **IDatabase cache=Connection.GetDatabase();**

B. IDatabase cache=Connection.GetCache();

✓ C. ICache cache=Connection.GetDatabase(); ✗

D. ICache cache=Connection.GetCache();

Explanation:

Answer – A

The right way is to use the IDatabase interface. Also you need to use the GetDatabase() method. This is also mentioned in the Microsoft documentation.

```
static void Main(string[] args)
{
    // Connection refers to a property that returns a ConnectionMultiplexer
    // as shown in the previous example.
    IDatabase cache = lazyConnection.Value.GetDatabase();
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on an example on how to work with Azure Redis from .Net, one can go to the below link

https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-dotnet-how-to-use-azure-redis-cache

Ask our Experts Rate this Question? ☺ ☹

View Queries open

Question 7 Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A development team is developing an application. The application will be working with customer data. The application will also be making use of Azure Redis Cache. You need to invalidate the cache when the customer data is changed.
You have to complete the below code to comply with the requirement

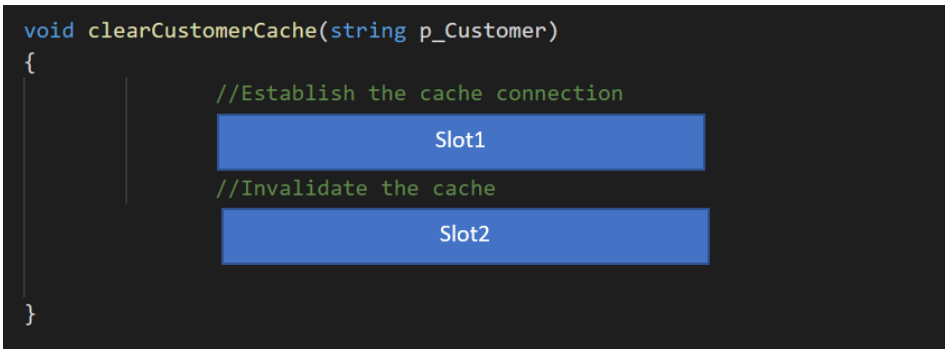

Which of the following will go into Slot2?

- A. **cache.KeyDelete(p_Customer);**
- B. **cache.ValueDelete(p_Customer);**
- C. **cache.StringGet(p_Customer);**
- D. **cache.StringSet(p_Customer);**

Explanation:

Answer – A

Since you have to invalidate the cache, you have to delete the Key itself

Option B is incorrect since you need to work with keys and not the values

Option C is incorrect this is used to get the string value

Option D is incorrect this is used to set the string value

For more information on an example on how to work with Azure Redis from .Net, one can go to the below link

https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-dotnet-how-to-use-azure-redis-cache

Ask our Experts Rate this Question?

View Queries open

Question 8 Correct

Domain :Develop Azure compute solutions

As a developer you need to create a Dockerfile for an application. The application will be based on ASP.Net core. The

application has the following requirements

Ensure that the application whizlabsApp.dll runs at the startup of the docker container

Run a powershell script called whizlabscsript.ps1 in the Docker container

The whizlabsApp.dll and the whizlabscsript.ps1 are in the same location as the DockerFile. Which of the following 3 coptions would be correct command, you place in the DockerFile?

- ✓ A. **FROM microsoft/dotnet:2.2-aspnetcore-runtime** ✓
- B. **EXPOSE whizlabsApp.dll ,whizlabscsript.ps1**
- ✓ C. **ENTRYPOINT ["dotnet", "whizlabsApp.dll"]** ✓
- D. **ENTRYPOINT ["whizlabsApp.dll" ,"whizlabscsript.ps1"]**
- ✓ E. **RUN powershell "whizlabscsript.ps1"** ✓
- F. **RUN "whizlabsApp.dll","whizlabscsript.ps1"**

Explanation:

Answer - A,C and E

Examples of DockerFiles are given in the Microsoft documentation. The below example shows how to define the base image and run an application on startup of the docker container

The following example shows a sample Dockerfile for an ASP.NET Core container.

Dockerfile Copy

```
FROM microsoft/dotnet:2.2-aspnetcore-runtime
ARG source
WORKDIR /app
EXPOSE 80
COPY ${source:-obj/Docker/publish} .
ENTRYPOINT ["dotnet", " MySingleContainerWebApp.dll "]
```

In this case, the image is based on version 2.2 of the official ASP.NET Core Docker image (multi-arch for Linux and Windows). This is the setting `FROM microsoft/dotnet:2.2-aspnetcore-runtime`. (For more information about this base image, see the .NET Core Docker Image page.) In the Dockerfile, you also need to instruct Docker to listen on the TCP port you will use at runtime (in this case, port 80, as configured with the EXPOSE setting).

You can specify additional configuration settings in the Dockerfile, depending on the language and framework you're using. For instance, the ENTRYPOINT line with `["dotnet", "MySingleContainerWebApp.dll"]` tells Docker to run a .NET Core application. If you're using the SDK and the .NET Core CLI (dotnet CLI) to build and run the .NET application, this setting would be different. The bottom line is that the ENTRYPOINT line and other settings will be different depending on the language and platform you choose for your application.

And the below example shows how to run a powershell script

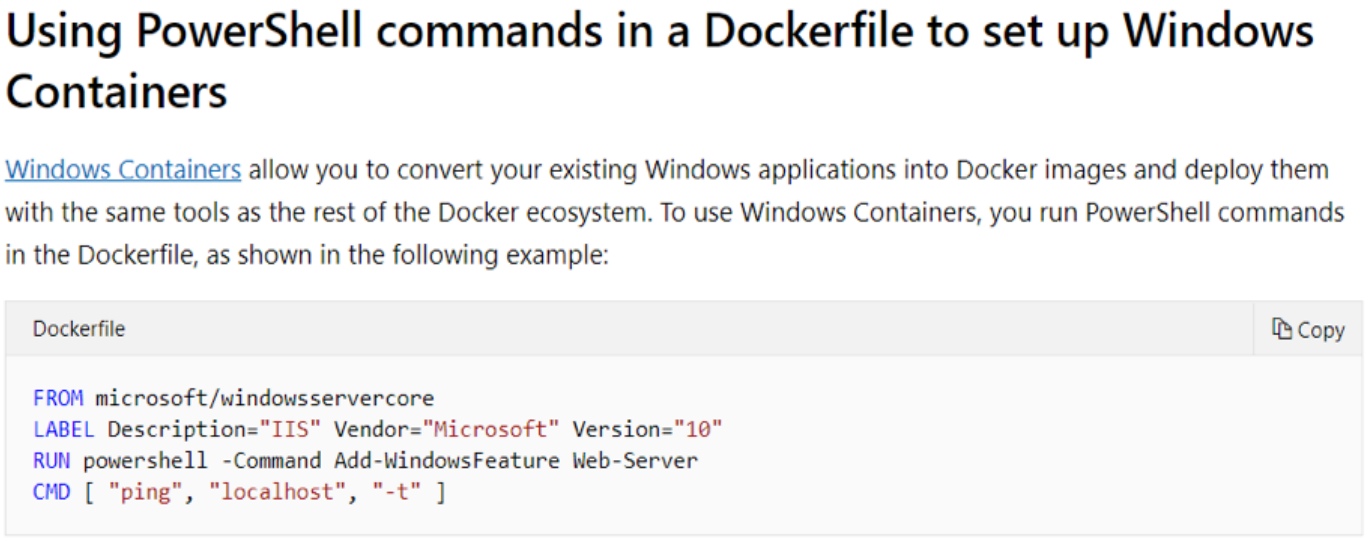

Using PowerShell commands in a Dockerfile to set up Windows Containers

Windows Containers allow you to convert your existing Windows applications into Docker images and deploy them with the same tools as the rest of the Docker ecosystem. To use Windows Containers, you run PowerShell commands in the Dockerfile, as shown in the following example:

Dockerfile Copy

```
FROM microsoft/windowsservercore
LABEL Description="IIS" Vendor="Microsoft" Version="10"
RUN powershell -Command Add-WindowsFeature Web-Server
CMD [ "ping", "localhost", "-t" ]
```

Based on the examples given in the documentation, all other options are incorrect

For more information on a complete docker application workflow, one can go to the below link

https://docs.microsoft.com/en-us/dotnet/standard/microservices-architecture/docker-application-development-process/docker-app-development-workflow

Ask our Experts Rate this Question?

View Queries open

Question 9 Correct

Domain :Connect to and consume Azure services and third-party services

A company is developing a solution that allows smart devices to send information to a central location. The solution must receive and store messages until they can be processed. You are planning to create an Azure Service Bus instance by providing a name, pricing tier, subscription, resource group, and location.
Which Azure CLI command or PowerShell cmdlet should you run?

✓ A. New-AzServiceBusQueue
-ResourceGroupName "whizlabs-rg"
-NamespaceName "whizlabs"
-Name "whizlabsqueue"
-EnablePartitioning $False

B. az group create --name "whizlabs-rg" --location "Central US"

C. New-AzResourceGroup -Name "whizlabs-rg" -Location "Central US"

D. New-AzServiceBusNamespace -ResourceGroup "whizlabs-rg" -NamespaceName "whizlabs" -Location "WestUS" -SkuName "Standard"

Explanation:

Answer – A

We need to create an Azure Service Buse Queue

The full syntax in powershell is given at

https://docs.microsoft.com/en-us/powershell/module/az.servicebus/new-azservicebusqueue?view=azps-4.4.0

New-AzServiceBusQueue
[-ResourceGroupName] <String>
[-Namespace] <String>
[-Name] <String>
[-EnablePartitioning <Boolean>]
[-LockDuration <String>]
[-AutoDeleteOnIdle <String>]
[-DefaultMessageTimeToLive <String>]
[-DuplicateDetectionHistoryTimeWindow <String>]
[-DeadLetteringOnMessageExpiration <Boolean>]
[-EnableBatchedOperations]
[-EnableExpress <Boolean>]
[-MaxDeliveryCount <Int32>]
[-MaxSizeInMegabytes <Int64>]
[-MessageCount <Int64>]
[-RequiresDuplicateDetection <Boolean>]
[-RequiresSession <Boolean>]
[-SizeInBytes <Int64>]
[-ForwardTo <String>]
[-ForwardDeadLetteredMessagesTo <String>]
[-DefaultProfile <IAzureContextContainer>]
[-WhatIf]
[-Confirm]
[<CommonParameters>]

Hence Option A is correct answer

Ask our Experts Rate this Question?

View Queries open

Question 10 Correct

Domain :Develop for Azure storage

A team needs to create an Azure CosmosDB account and ensure that regional failover's are in place for the account. Which of the following is a valid CLI command for creation of the CosmosDB account?

✓ A.
```
az cosmosdb update \
        --name "whizlabscosmosdb" \
        --resource-group "whizlabs-rg" \
        --locations "South Central US"=0 "North Central US"=1 "East US"=2
```

B.
```
az cosmosdb create \
        --name "whizlabscosmosdb" \
        --resource-group "whizlabs-rg" \
        --locations 3
```

C.
```
az cosmosdb create \
        --name "whizlabscosmosdb" \
        --resource-group "whizlabs-rg" \
        --locations "South Central US"=3
```

D.
```
az cosmosdb update \
        --name "whizlcreateabscosmosdb" \
        --resource-group "whizlabs-rg" \
        --enable-locations
```

Explanation:

Answer – A

The Microsoft documentation mentions the correct CLI command to use for ensuring a Cosmos DB account is created in multiple regions

Since this is clearly mentioned, all other options are incorrect

For more information on how to enable multi-region for CosmosDB, one can go to the below link

https://docs.microsoft.com/en-us/cli/azure/cosmosdb?view=azure-cli-latest#az-cosmosdb-create

Ask our Experts Rate this Question?

View Queries open

Question 11 Correct

Domain :Connect to and consume Azure services and third-party services

A developer has setup a web application in Azure and also setup Azure CDN to route requests to the Web App. One of the requirements is to ensure that if users make requests based on passing an ID parameter, then those requests should always be served from a Point of Presence. An example of the URL is given below
https://whizlabs.com/Customer.aspx?ID=1

Which of the following mode should be set for the query string setting for the CDN service?

A. Ignore query strings

B. Default setting

C. Bypass caching

✓ D. Cache every unique URL

Explanation:

Answer - D

Below are the different settings available for the CDN when it comes to caching of the query string.

Since we need to ensure that query strings are cached, we have to choose the option of 'Cache every unique URL'

- **Ignore query strings**: Default mode. In this mode, the CDN point-of-presence (POP) node passes the query strings from the requestor to the origin server on the first request and caches the asset. All subsequent requests for the asset that are served from the POP ignore the query strings until the cached asset expires.
- **Bypass caching for query strings**: In this mode, requests with query strings are not cached at the CDN POP node. The POP node retrieves the asset directly from the origin server and passes it to the requestor with each request.
- **Cache every unique URL**: In this mode, each request with a unique URL, including the query string, is treated as a unique asset with its own cache. For example, the response from the origin server for a request for example.ashx?q=test1 is cached at the POP node and returned for subsequent caches with the same query string. A request for example.ashx?q=test2 is cached as a separate asset with its own time-to-live setting.

Since this is the ideal approach, all other options are incorrect

For more information on working with query strings for CDN, one can go to the below link

https://docs.microsoft.com/en-us/azure/cdn/cdn-query-string

Ask our Experts Rate this Question?

View Queries open

Question 12 Correct

Domain :Develop Azure compute solutions

A developer needs to run a set of Azure CLI commands to create a virtual machine. You need to complete the below set of commands.

Which of the following would go into Slot1?

A. vm create

✓ B. group create

C. vm set

D. group set

Explanation:

Answer – B

An example of this is given in the Microsoft documentation. The first step is to ensure that you create a resource group via the

az group create CLI command

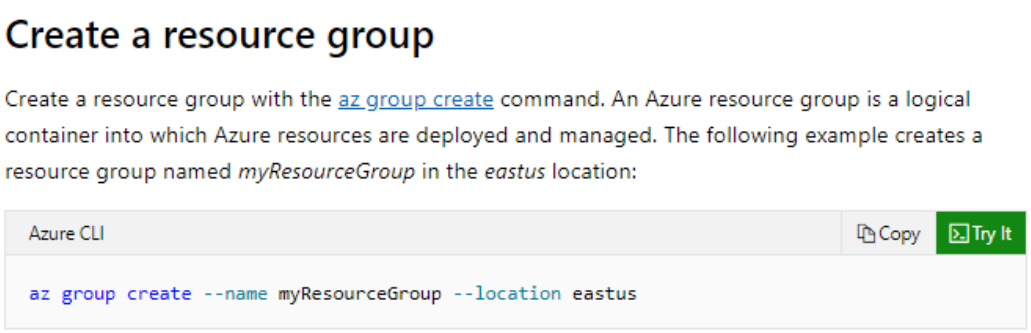

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on using Azure CLI commands to create virtual machines, please visit the below URL

https://docs.microsoft.com/en-us/azure/virtual-machines/windows/quick-create-cli

Ask our Experts Rate this Question?

View Queries open

Question 13 Correct

Domain :Develop Azure compute solutions

A developer needs to run a set of Azure CLI commands to create a virtual machine. You need to complete the below set of commands.

Which of the following would go into Slot2?

- ✓ A. vm create
- B. group create
- C. vm set
- D. group set

Explanation:

Answer – A

An example of this is given in the Microsoft documentation. The next step is to ensure that you create the virtual machine via the vm create CLI command

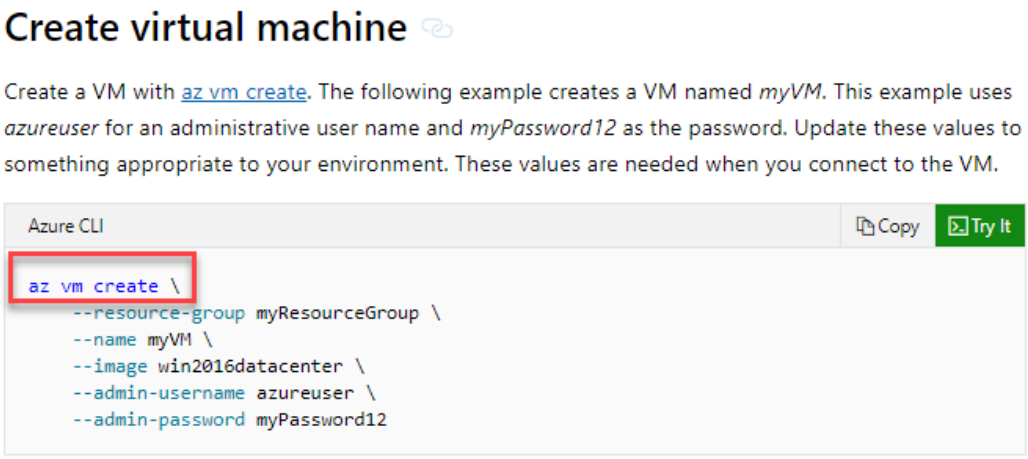

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on using Azure CLI commands to create virtual machines, please visit the below URL

https://docs.microsoft.com/en-us/azure/virtual-machines/windows/quick-create-cli

Ask our Experts Rate this Question?

View Queries open

Question 14 Correct

Domain :Develop Azure compute solutions

A developer needs to run a set of Azure CLI commands to create a virtual machine. You need to complete the below set of commands.

az Slot1 --name whizlabgroup --location eastus

az Slot2 \
--resource-group whizlabgroup \
--name whizlabvm \
Slot3 win2016datacenter \
--admin-username whizlabsr \
--admin-password Whizlabsr123

Which of the following would go into Slot3?

- A. --image
- B. --edition
- C. --version
- D. --create

Explanation:

Answer – A

Here, we are mentioning that the image name should be Windows Server 2016. An example of this is provided in the Microsoft documentation

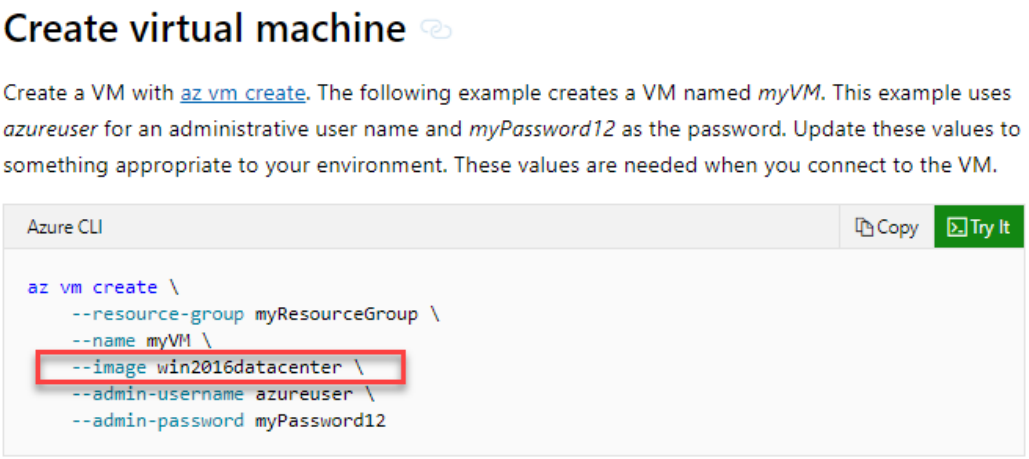

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on using Azure CLI commands to create virtual machines, please visit the below URL

https://docs.microsoft.com/en-us/azure/virtual-machines/windows/quick-create-cli

Ask our Experts Rate this Question?

View Queries open

Question 15 Correct

Domain :Develop Azure compute solutions

Your company has a set of Azure storage accounts. These storage accounts store blob objects. You have to move the blobs from one container to another across storage accounts.
You have to use the right tool to perform the movement of data.
You decide to use the AzCopy tool to implement this requirement
Would this fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – A

Yes, you can use the AzCopy tool for this requirement

The Microsoft documentation mentions the following

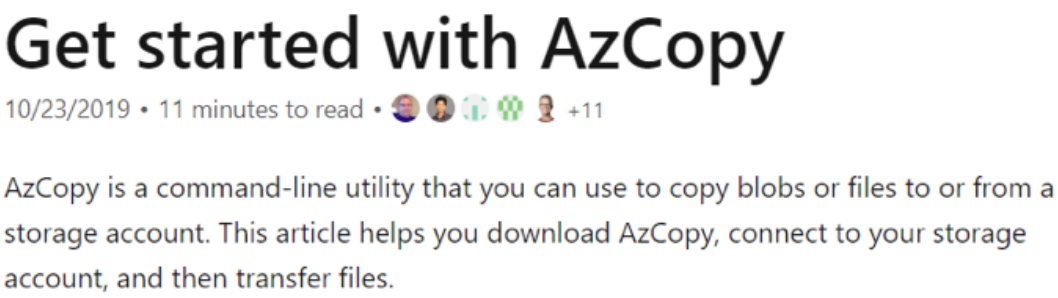

For more information on the AzCopy tool, please refer to the below link

https://docs.microsoft.com/en-us/azure/storage/common/storage-use-azcopy-v10

Ask our Experts Rate this Question?

View Queries open

Question 16 Correct

Domain :Develop Azure compute solutions

Your company has a set of Azure storage accounts. These storage accounts store blob objects. You have to move the blobs from one container to another across storage accounts.
You have to use the right tool to perform the movement of data.
You decide to use the Azure CLI tool to implement this requirement
Would this fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – A

Yes, you can also use the Azure CLI tool for this requirement

The below command can be used for this requirement

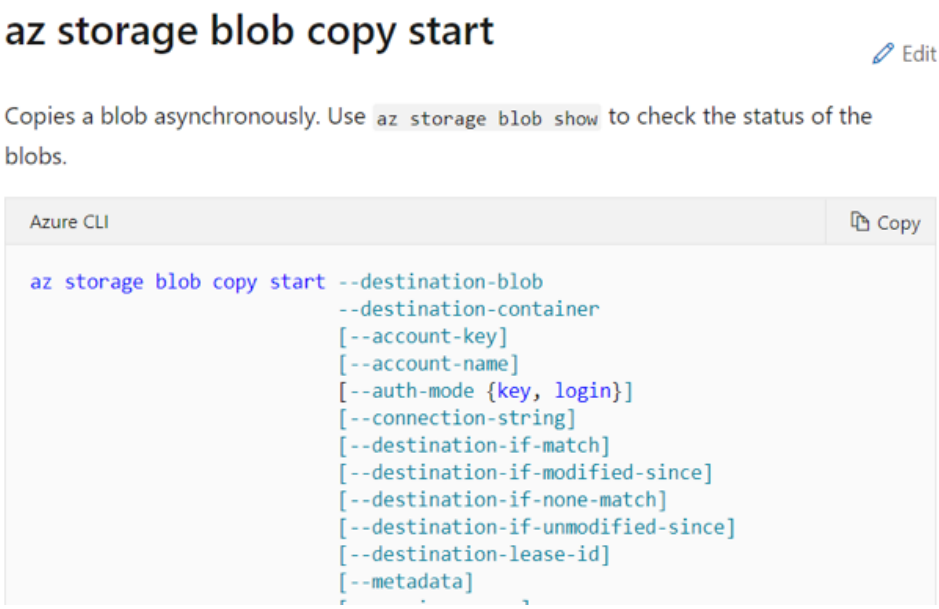

For more information on the Azure CLI command, please refer to the below link

https://docs.microsoft.com/en-us/cli/azure/storage/blob/copy?view=azure-cli-latest

Ask our Experts Rate this Question?

View Queries open

Question 17 Correct

Domain :Develop Azure compute solutions

Your company has a set of Azure storage accounts. These storage accounts store blob objects. You have to move the blobs from one container to another across storage accounts.
You have to use the right tool to perform the movement of data.
You decide to use the PowerShell tool to implement this requirement
Would this fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – A

Yes, you can also use the PowerShell tool for this requirement

The below command can be used for this requirement

Start-AzureStorageBlobCopy

Module: Azure.Storage

Starts to copy a blob.

PowerShell Copy

```
Start-AzureStorageBlobCopy
     [-SrcBlob] <String>
     -SrcContainer <String>
     -DestContainer <String>
     [-DestBlob <String>]
     [-PremiumPageBlobTier <PremiumPageBlobTier>]
     [-Context <IStorageContext>]
     [-DestContext <IStorageContext>]
     [-Force]
     [-ServerTimeoutPerRequest <Int32>]
     [-ClientTimeoutPerRequest <Int32>]
     [-DefaultProfile <IAzureContextContainer>]
     [-ConcurrentTaskCount <Int32>]
     [-WhatIf]
     [-Confirm]
     [<CommonParameters>]
```

For more information on the PowerShell command, please refer to the below link

https://docs.microsoft.com/en-us/powershell/module/azure.storage/start-azurestorageblobcopy

Ask our Experts Rate this Question?

View Queries open

Question 18 Correct

Domain :Develop Azure compute solutions

Your team has an Azure container registry in place. You need to provide a set of developers the ability to publish images to registry. You need to ensure the least privilege access is given to the developers. Which of the following role would you give for this purpose?

A. Owner

B. Contributor

✓ C. AcrPush

D. AcrPull

Explanation:

Answer – C

You would provide the AcrPush role. The different roles are given in the Microsoft documentation

Role/Permission	Access Resource Manager	Create/delete registry	Push image	Pull image	Delete image data	Change policies	Sign images
Owner	X	X	X	X	X	X	
Contributor	X	X	X	X	X	X	
Reader	X			X			
AcrPush			X	X			
AcrPull				X			
AcrDelete					X		
AcrImageSigner							X

Options A and B are incorrect because these roles would give more access than what is required.

Option D is incorrect because this does not provide the ability to push images to the registry.

For more information on Azure container registry roles, please visit the following URL

https://docs.microsoft.com/en-us/azure/container-registry/container-registry-roles

Ask our Experts

Rate this Question?

View Queries open

Question 19 Correct

Domain :Develop for Azure storage

A company is currently looking at using Azure storage accounts as a data storage platform. They want to be able to store documents that would need to be retained for around a year. They want to minimize the cost of storage. These documents are not accessed that frequently.
Which of the following tier would you use for storing the files?

A. Hot

B. Cool

✓ C. Archive

D. Secondary

Explanation:

Answer – C

In the BLOB service, you can use the Archive tier to store archive documents that can save on costs. The Microsoft documentation mentions the following

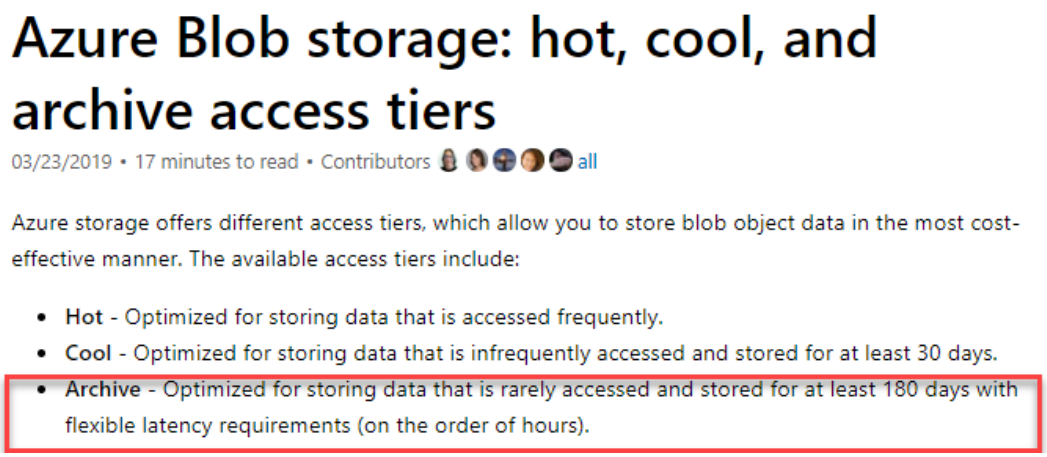

Azure Blob storage: hot, cool, and archive access tiers

03/23/2019 • 17 minutes to read • Contributors all

Azure storage offers different access tiers, which allow you to store blob object data in the most cost-effective manner. The available access tiers include:

- **Hot** - Optimized for storing data that is accessed frequently.
- **Cool** - Optimized for storing data that is infrequently accessed and stored for at least 30 days.
- **Archive** - Optimized for storing data that is rarely accessed and stored for at least 180 days with flexible latency requirements (on the order of hours).

Options A and B are incorrect since these are not cost-efficient options for storing archive-based files.

Option D is incorrect since this is not a valid based tier.

For more information on blob storage tiers, please visit the below URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-storage-tiers

Ask our Experts Rate this Question?

View Queries open

Question 20 Correct

Domain :Develop for Azure storage

A company is currently looking at using Azure storage accounts as a data storage platform. They want to be able to store documents that would need to be retained for around a year. They want to minimize the cost of storage. These documents are not accessed that frequently.
If you need to retrieve an object in Azure BLOB Archive storage, which of the following would need to be done first?

A. **Change the access permissions**

✓ B. **Change the tier of the object**

C. **Change the account kind**

D. **Change the access keys**

Explanation:

Answer – B

You need to change the tier of the object. This is also mentioned in the Microsoft documentation

Blob rehydration

To read data in archive storage, you must first change the tier of the blob to hot or cool. This process is known as rehydration and can take up to 15 hours to complete. Large blob sizes are recommended for optimal performance. Rehydrating several small blobs concurrently may add additional time.

During rehydration, you may check the **Archive Status** blob property to confirm if the tier has changed. The status reads "rehydrate-pending-to-hot" or "rehydrate-pending-to-cool" depending on the destination tier. Upon completion, the archive status property is removed, and the **Access Tier** blob property reflects the new hot or cool tier.

Since this is clearly mentioned in the documentation, all other options are incorrect

For more information on blob storage tiers, please visit the below URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-storage-tiers

Ask our Experts Rate this Question?

View Queries open

Question 21 Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A developer has been assigned a task to create code which would interact with an Azure Redis instance. Objects of the following class need to be uploaded to the Azure Redis Cache database

```
class WhizlabCustomer
        {
            public string Id { get; set; }
            public string Name { get; set; }

            public Employee(string pID, string pName)
            {
                this.Id = pID;
                this.Name = pName;
            }
        }
```

You need to complete the below code snippet

```
// Code to store the object in cache
        WhizlabCustomer obj = new WhizlabCustomer("1", "David");

        cache. Slot1 ("ID1", JsonConvert. Slot2 (obj)));

// Retrieve the object from the cache
WhizlabCustomer objcache = JsonConvert. Slot3 <WhizlabCustomer>(cache. Slot4 ("ID1"));
```

Which of the following would go into Slot1?

A. **ObjectSet**

✓ B. **StringSet**

C. **ClassSet**

D. **Set**

Explanation:

Answer – B

To Add an object to the cache database, we need to use the StringSet Method. An example of this is also given in the Microsoft documentation

Add the following `Employee` class definition to *Program.cs*:

C# Copy

```
class Employee
{
    public string Id { get; set; }
    public string Name { get; set; }
    public int Age { get; set; }

    public Employee(string EmployeeId, string Name, int Age)
    {
        this.Id = EmployeeId;
        this.Name = Name;
        this.Age = Age;
    }
}
```

At the bottom of `Main()` procedure in *Program.cs*, and before the call to `Dispose()`, add the following lines of code to cache and retrieve a serialized .NET object:

C# Copy

```
// Store .NET object to cache
Employee e007 = new Employee("007", "Davide Columbo", 100);
Console.WriteLine("Cache response from storing Employee .NET object : " +
    cache.StringSet("e007", JsonConvert.SerializeObject(e007)));

// Retrieve .NET object from cache
Employee e007FromCache = JsonConvert.DeserializeObject<Employee>(cache.StringGet("e007"));
Console.WriteLine("Deserialized Employee .NET object :\n");
Console.WriteLine("\tEmployee.Name : " + e007FromCache.Name);
Console.WriteLine("\tEmployee.Id   : " + e007FromCache.Id);
Console.WriteLine("\tEmployee.Age  : " + e007FromCache.Age + "\n");
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on how to work with Azure Redis Cache from a .Net program, please visit the below URL

https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-dotnet-how-to-use-azure-redis-cache

Ask our Experts Rate this Question?

View Queries open

Question 22 Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A developer has been assigned a task to create code which would interact with an Azure Redis instance. Objects of the

following class need to be uploaded to the Azure Redis Cache database

```
class WhizlabCustomer
        {
            public string Id { get; set; }
            public string Name { get; set; }

            public Employee(string pID, string pName)
            {
                this.Id = pID;
                this.Name = pName;
            }
        }
```

You need to complete the below code snippet

Which of the following would go into Slot2?

A. SetObject

✓ B. SerializeObject

C. GetObject

D. SerializeClass

Explanation:

Answer – B

We need to serialize the object before it can be added to the redis cache database.

An example of this is also given in the Microsoft documentation

Add the following `Employee` class definition to *Program.cs*:

C# Copy

```
class Employee
{
    public string Id { get; set; }
    public string Name { get; set; }
    public int Age { get; set; }

    public Employee(string EmployeeId, string Name, int Age)
    {
        this.Id = EmployeeId;
        this.Name = Name;
        this.Age = Age;
    }
}
```

At the bottom of `Main()` procedure in *Program.cs*, and before the call to `Dispose()`, add the following lines of code to cache and retrieve a serialized .NET object:

C# Copy

```
// Store .NET object to cache
Employee e007 = new Employee("007", "Davide Columbo", 100);
Console.WriteLine("Cache response from storing Employee .NET object : " +
    cache.StringSet("e007", JsonConvert.SerializeObject(e007)));

// Retrieve .NET object from cache
Employee e007FromCache = JsonConvert.DeserializeObject<Employee>(cache.StringGet("e007"));
Console.WriteLine("Deserialized Employee .NET object :\n");
Console.WriteLine("\tEmployee.Name : " + e007FromCache.Name);
Console.WriteLine("\tEmployee.Id   : " + e007FromCache.Id);
Console.WriteLine("\tEmployee.Age  : " + e007FromCache.Age + "\n");
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on how to work with Azure Redis Cache from a .Net program, please visit the below URL

https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-dotnet-how-to-use-azure-redis-cache

Ask our Experts Rate this Question?

View Queries open

Question 23 Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A developer has been assigned a task to create code which would interact with an Azure Redis instance. Objects of the

following class need to be uploaded to the Azure Redis Cache database

```
class WhizlabCustomer
        {
            public string Id { get; set; }
            public string Name { get; set; }

            public Employee(string pID, string pName)
            {
                this.Id = pID;
                this.Name = pName;
            }
        }
```

You need to complete the below code snippet

```
// Code to store the object in cache
        WhizlabCustomer obj = new WhizlabCustomer("1", "David");

        cache. Slot1 ("ID1", JsonConvert. Slot2 (obj)));

// Retrieve the object from the cache
WhizlabCustomer objcache = JsonConvert. Slot3 <WhizlabCustomer>(cache. Slot4 ("ID1"));
```

Which of the following would go into Slot3?

A. **SetObject**

✓ B. **DeserializeObject**

C. **GetObject**

D. **DeserializeClass**

Explanation:

Answer – B

We need to use the deserialize method to convert the object retrieved from the cache

An example of this is also given in the Microsoft documentation

Add the following `Employee` class definition to *Program.cs*:

C# Copy

```
class Employee
{
    public string Id { get; set; }
    public string Name { get; set; }
    public int Age { get; set; }

    public Employee(string EmployeeId, string Name, int Age)
    {
        this.Id = EmployeeId;
        this.Name = Name;
        this.Age = Age;
    }
}
```

At the bottom of `Main()` procedure in *Program.cs*, and before the call to `Dispose()`, add the following lines of code to cache and retrieve a serialized .NET object:

C# Copy

```
// Store .NET object to cache
Employee e007 = new Employee("007", "Davide Columbo", 100);
Console.WriteLine("Cache response from storing Employee .NET object : " +
    cache.StringSet("e007", JsonConvert.SerializeObject(e007)));

// Retrieve .NET object from cache
Employee e007FromCache = JsonConvert.DeserializeObject<Employee>(cache.StringGet("e007"));
Console.WriteLine("Deserialized Employee .NET object :\n");
Console.WriteLine("\tEmployee.Name : " + e007FromCache.Name);
Console.WriteLine("\tEmployee.Id   : " + e007FromCache.Id);
Console.WriteLine("\tEmployee.Age  : " + e007FromCache.Age + "\n");
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on how to work with Azure Redis Cache from a .Net program, please visit the below URL

https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-dotnet-how-to-use-azure-redis-cache

Ask our Experts Rate this Question?

View Queries open

Question 24 Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A developer has been assigned a task to create code which would interact with an Azure Redis instance. Objects of the

following class need to be uploaded to the Azure Redis Cache database

```
class WhizlabCustomer
        {
            public string Id { get; set; }
            public string Name { get; set; }

            public Employee(string pID, string pName)
            {
                this.Id = pID;
                this.Name = pName;
            }
        }
```

You need to complete the below code snippet

```
// Code to store the object in cache
        WhizlabCustomer obj = new WhizlabCustomer("1", "David");

        cache. Slot1 ("ID1", JsonConvert. Slot2 (obj)));

// Retrieve the object from the cache
WhizlabCustomer objcache = JsonConvert. Slot3 <WhizlabCustomer>(cache. Slot4 ("ID1"));
```

Which of the following would go into Slot4?

A. **ObjectGet**

B. **StringSet**

C. **ClassSet**

✓ D. **StringGet**

Explanation:

Answer – D

We need to get an object from the cache, so we should use the StringGet method.

An example of this is also given in the Microsoft documentation

Add the following `Employee` class definition to *Program.cs*:

C# Copy

```
class Employee
{
    public string Id { get; set; }
    public string Name { get; set; }
    public int Age { get; set; }

    public Employee(string EmployeeId, string Name, int Age)
    {
        this.Id = EmployeeId;
        this.Name = Name;
        this.Age = Age;
    }
}
```

At the bottom of `Main()` procedure in *Program.cs*, and before the call to `Dispose()`, add the following lines of code to cache and retrieve a serialized .NET object:

C# Copy

```
// Store .NET object to cache
Employee e007 = new Employee("007", "Davide Columbo", 100);
Console.WriteLine("Cache response from storing Employee .NET object : " +
    cache.StringSet("e007", JsonConvert.SerializeObject(e007)));

// Retrieve .NET object from cache
Employee e007FromCache = JsonConvert.DeserializeObject<Employee>(cache.StringGet("e007"));
Console.WriteLine("Deserialized Employee .NET object :\n");
Console.WriteLine("\tEmployee.Name : " + e007FromCache.Name);
Console.WriteLine("\tEmployee.Id   : " + e007FromCache.Id);
Console.WriteLine("\tEmployee.Age  : " + e007FromCache.Age + "\n");
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on how to work with Azure Redis Cache from a .Net program, please visit the below URL

https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-dotnet-how-to-use-azure-redis-cache

Ask our Experts Rate this Question?

View Queries open

Question 25 Correct

Domain :Develop for Azure storage

A company is developing a system which is going to be using Azure Cosmos DB at the underlying data store. Below are the

requirements of the data store

Ensure at least 99.99% availability and provide network failures

Accepts writes via the application even in the case of network outages or any unforeseen failures

Process data in the same sequence as the writes being made

Allow out of order data with a maximum of 5 second tolerance window

You have to provision a Cosmos DB account – SQL API. You already have a resource group in the South Central US region. You have to complete the below Azure CLI commands for this purpose.

```
resourceGroup='whizlabs-rg'
accountname='whizlabacc'
databasename='whizlabdb'
collectionName='whizlabcollection'
consistencyLevel='  Slot1  '
az cosmosdb create --name $accountname \
  Slot2
--resource-group $ resourceGroup –max-interval 5 \
--default-consistency-level=$ consistencyLevel
  Slot3
```

Which of the following would go into Slot1?

A. **Strong**

B. **Eventual**

C. **ConsistentPrefix**

✓ D. **BoundedStaleness**

Explanation:

Answer – D

Since you can have an out of order read for a maximum of 5 seconds, this becomes our staleness window.

The Microsoft documentation mentions the following on the Bounded Staleness consistency level.

- **Bounded staleness**: The reads are guaranteed to honor the consistent-prefix guarantee. The reads might lag behind writes by at most "*K*" versions (i.e., "updates") of an item or by "*T*" time interval. In other words, when you choose bounded staleness, the "staleness" can be configured in two ways:
  - The number of versions (*K*) of the item
  - The time interval (*T*) by which the reads might lag behind the writes

  Bounded staleness offers total global order except within the "staleness window." The monotonic read guarantees exist within a region both inside and outside the staleness window. Strong consistency has the same semantics as the one offered by bounded staleness. The staleness window is equal to zero. Bounded staleness is also referred to as time-delayed linearizability. When a client performs read operations within a region that accepts writes, the guarantees provided by bounded staleness consistency are identical to those guarantees by the strong consistency.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on consistency levels, please visit the below URL

https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels

Ask our Experts Rate this Question?

View Queries open

Question 26 Correct

Domain :Develop for Azure storage

A company is developing a system which is going to be using Azure Cosmos DB at the underlying data store. Below are the requirements of the data store

Ensure at least 99.99% availability and provide network failures

Accepts writes via the application even in the case of network outages or any unforeseen failures

Process data in the same sequence as the writes being made

Allow out of order data with a maximum of 5 second tolerance window

You have to provision a Cosmos DB account – SQL API. You already have a resource group in the South Central US region. You have to complete the below Azure CLI commands for this purpose.

```
resourceGroup='whizlabs-rg'
accountname='whizlabacc'
databasename='whizlabdb'
collectionName='whizlabcollection'
consistencyLevel='  Slot1  '
az cosmosdb create --name $accountname \
  Slot2
--resource-group $ resourceGroup –max-interval 5 \
--default-consistency-level=$ consistencyLevel
  Slot3
```

Which of the following would go into Slot2?

A. **--enable-virtual-network true**

✓ B. **--enable-automatic-failover true**

C. **--kind 'GlobalDocumentDB'**

D. **--kind 'MongoDB'**

Explanation:

Answer – B

Since we have to ensure that the data needs to be available even in the case of network outages or any unforeseen failures, we have to enable automatic failover.

The Microsoft documentation mentions the following

`--enable-automatic-failover`

Enables automatic failover of the write region in the rare event that the region is unavailable due to an outage. Automatic failover will result in a new write region for the account and is chosen based on the failover priorities configured for the account.
accepted values: false, true

Option A is incorrect since there is no mention in the question of requiring the database to be part of a virtual network

Option C is incorrect since the default API chosen for the database is the SQL API

Option D is incorrect since we need to create a Cosmos DB account with the SQL API

For more information on the Cosmos DB create command, please visit the below URL

https://docs.microsoft.com/en-us/cli/azure/cosmosdb?view=azure-cli-latest#az-cosmosdb-create

Ask our Experts Rate this Question?

View Queries open

Question 27 Correct

Domain :Develop for Azure storage

A company is developing a system which is going to be using Azure Cosmos DB at the underlying data store. Below are the requirements of the data store

Ensure at least 99.99% availability and provide network failures

Accepts writes via the application even in the case of network outages or any unforeseen failures

Process data in the same sequence as the writes being made

Allow out of order data with a maximum of 5 second tolerance window

You have to provision a Cosmos DB account – SQL API. You already have a resource group in the South Central US region. You have to complete the below Azure CLI commands for this purpose.

```
resourceGroup='whizlabs-rg'
accountname='whizlabacc'
databasename='whizlabdb'
collectionName='whizlabcollection'
consistencyLevel='  Slot1  '
az cosmosdb create --name $accountname \
  Slot2
--resource-group $ resourceGroup –max-interval 5 \
--default-consistency-level=$ consistencyLevel
  Slot3
```

Which of the following would go into Slot3?

A. --locations 'southecentralus'

B. --locations 'eastus'

✓ C. --locations 'southcentalus=0 eastus=1'

D. --locations 'southecentralus=0'

Explanation:

Answer – C

Since we need to have additional regions for failover purpose, we need to add multiple locations to the Cosmos DB account.

The other options are incorrect since they only have one location specified.

For more information on the Cosmos DB create command, please visit the below URL

https://docs.microsoft.com/en-us/cli/azure/cosmosdb?view=azure-cli-latest#az-cosmosdb-create

Ask our Experts

Rate this Question?

View Queries open

Question 28

Correct

Domain :Develop for Azure storage

You have to implement the azcopy tool to copy objects from a local folder named D:\whizlabs to a container named "demo" within the below storage account

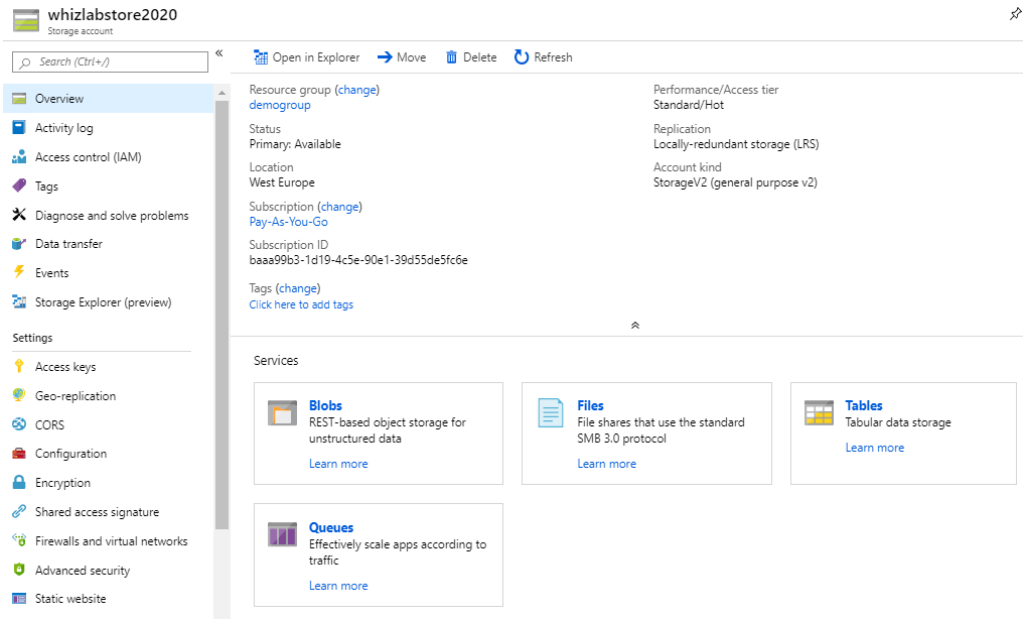

You have to complete the below command to copy all of the objects in the local folder

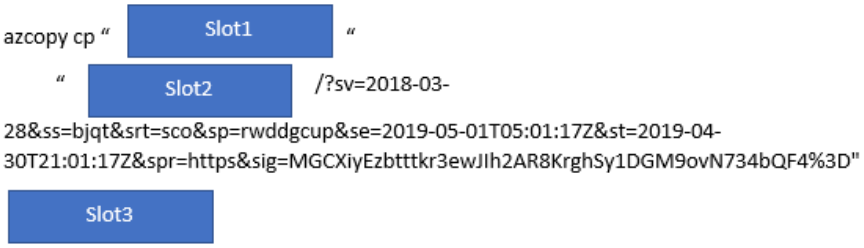

Which of the following would go into Slot1?

A. https://whizlabstore2020.blob.core.windows.net/demo

B. https://whizlabstore2020/demo

✓ C. D:\whizlabs

D. whizlabs

Explanation:

Answer – C

Since we need to copy objects from the local folder, we have to mention the entire local folder path.

An example of this is given in the Microsoft documentation

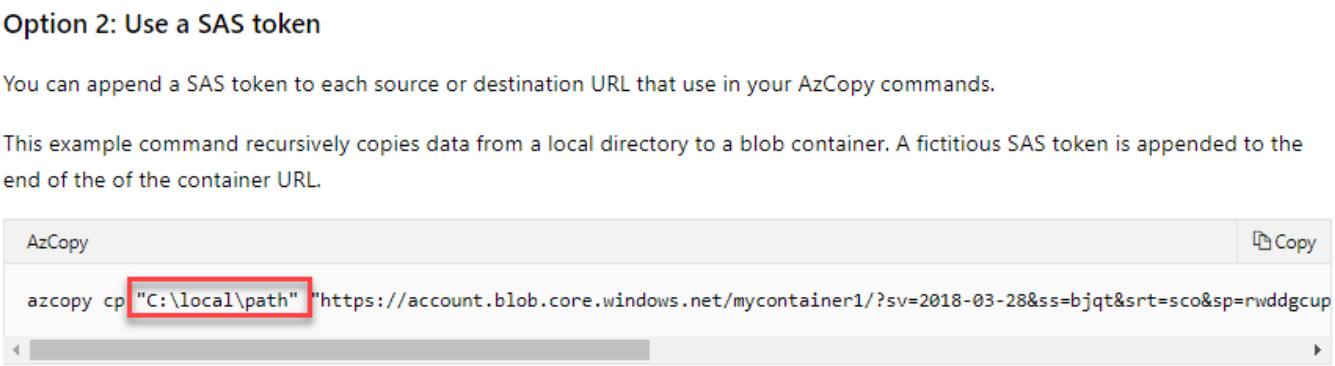

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on using the AzCopy tool, please visit the below URL

https://docs.microsoft.com/en-us/azure/storage/common/storage-use-azcopy-v10

Ask our Experts Rate this Question?

View Queries open

Question 29 Correct

Domain :Develop for Azure storage

You have to implement the azcopy tool to copy objects from a local folder named D:\whizlabs to a container named "demo" within the below storage account

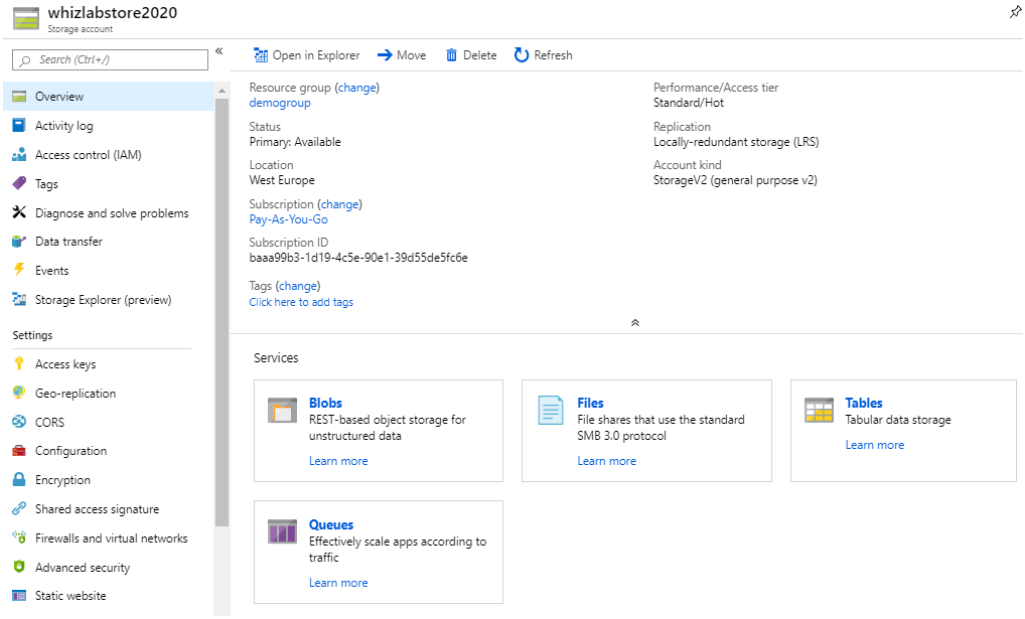

You have to complete the below command to copy all of the objects in the local folder

azcopy cp " Slot1 "

" Slot2 /?sv=2018-03-28&ss=bjqt&srt=sco&sp=rwddgcup&se=2019-05-01T05:01:17Z&st=2019-04-30T21:01:17Z&spr=https&sig=MGCXiyEzbtttkr3ewJIh2AR8KrghSy1DGM9ovN734bQF4%3D"

Slot3

Which of the following would go into Slot2?

- ✓ A. https://whizlabstore2020.blob.core.windows.net/demo
- B. https://whizlabstore2020/demo
- C. D:\whizlabs
- D. whizlabs

Explanation:

Answer – A

Here since we need to copy it to the container, we have to mention the full URI of the container

An example of this is given in the Microsoft documentation

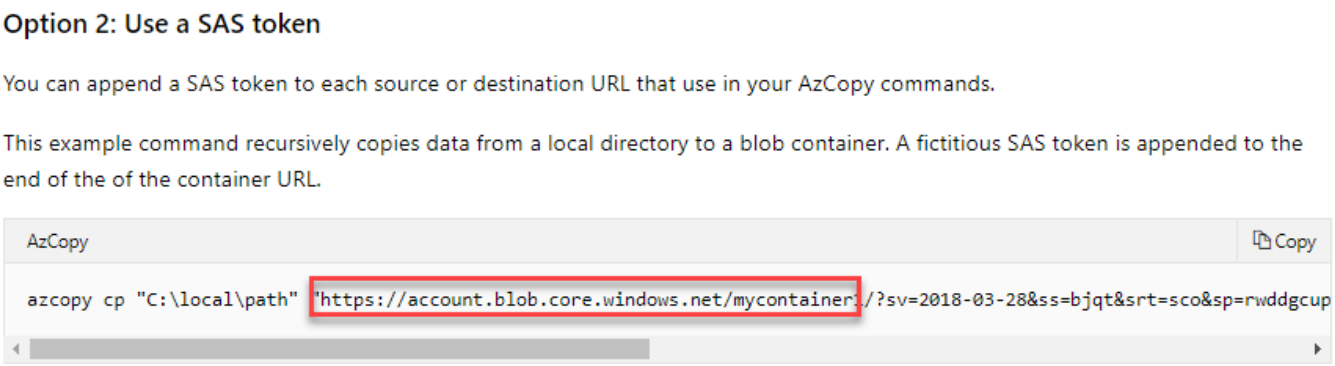

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on using the AzCopy tool, please visit the below URL

https://docs.microsoft.com/en-us/azure/storage/common/storage-use-azcopy-v10

Ask our Experts Rate this Question?

View Queries open

Question 30 Correct

Domain :Develop for Azure storage

You have to implement the azcopy tool to copy objects from a local folder named D:\whizlabs to a container named "demo" within the below storage account

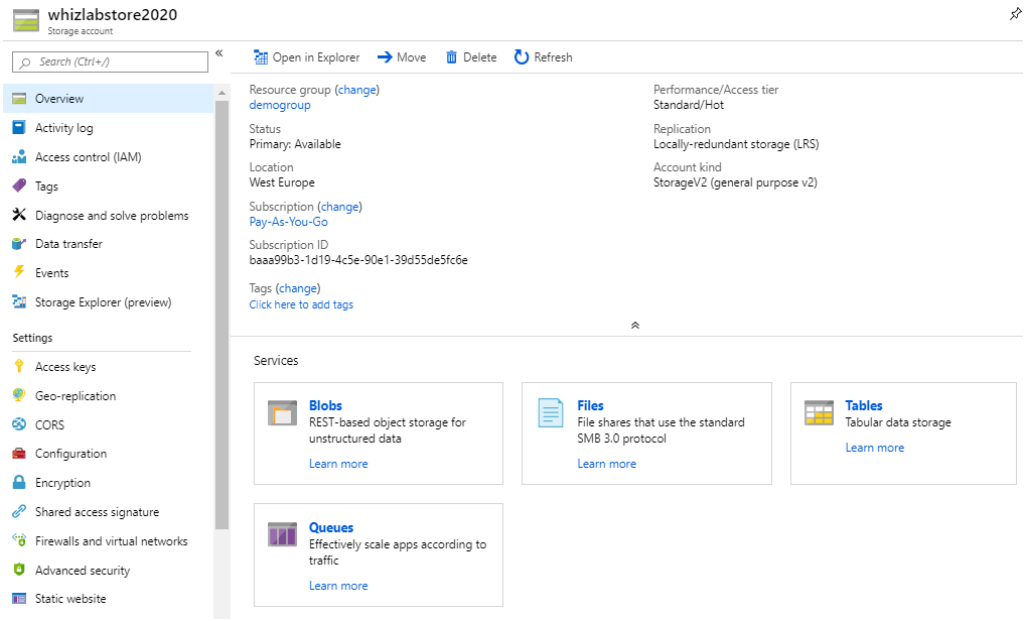

You have to complete the below command to copy all of the objects in the local folder

azcopy cp " Slot1 "

" Slot2 /?sv=2018-03-28&ss=bjqt&srt=sco&sp=rwddgcup&se=2019-05-01T05:01:17Z&st=2019-04-30T21:01:17Z&spr=https&sig=MGCXiyEzbtttkr3ewJIh2AR8KrghSy1DGM9ovN734bQF4%3D"

Slot3

Which of the following would go into Slot3?

A. --copy-all

✓ B. --recursive=true

C. --tree=true

D. --copy-files=true

Explanation:

Answer – B

Here the option for copying all the files recursively is --recursive=true

An example of this is given in the Microsoft documentation

Option 2: Use a SAS token

You can append a SAS token to each source or destination URL that use in your AzCopy commands.

This example command recursively copies data from a local directory to a blob container. A fictitious SAS token is appended to the end of the of the container URL.

AzCopy Copy

```
9-05-01T05:01:17Z&st=2019-04-30T21:01:17Z&spr=https&sig=MGCXiyEzbtttkr3ewJIh2AR8KrghSy1DGM9ovN734bQF4%3D" --recursive=true
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on using the AzCopy tool, please visit the below URL

https://docs.microsoft.com/en-us/azure/storage/common/storage-use-azcopy-v10

Ask our Experts Rate this Question?

View Queries open

Question 31 Correct

Domain :Develop for Azure storage

View Case Study

You need to formulate a query that would be used to get all items from container customer where the order price is $9.99

```
select * from c IN  Slot1  where  Slot2  ="9.99"
```

Which of the following would go into Slot1?

A. orders

B. customer

C. course

✓ D. customer.orders

Explanation:

Answer – D

We can use the IN clause to query data from JSON arrays

An example of an implementation of querying the data is given below

Data in the container

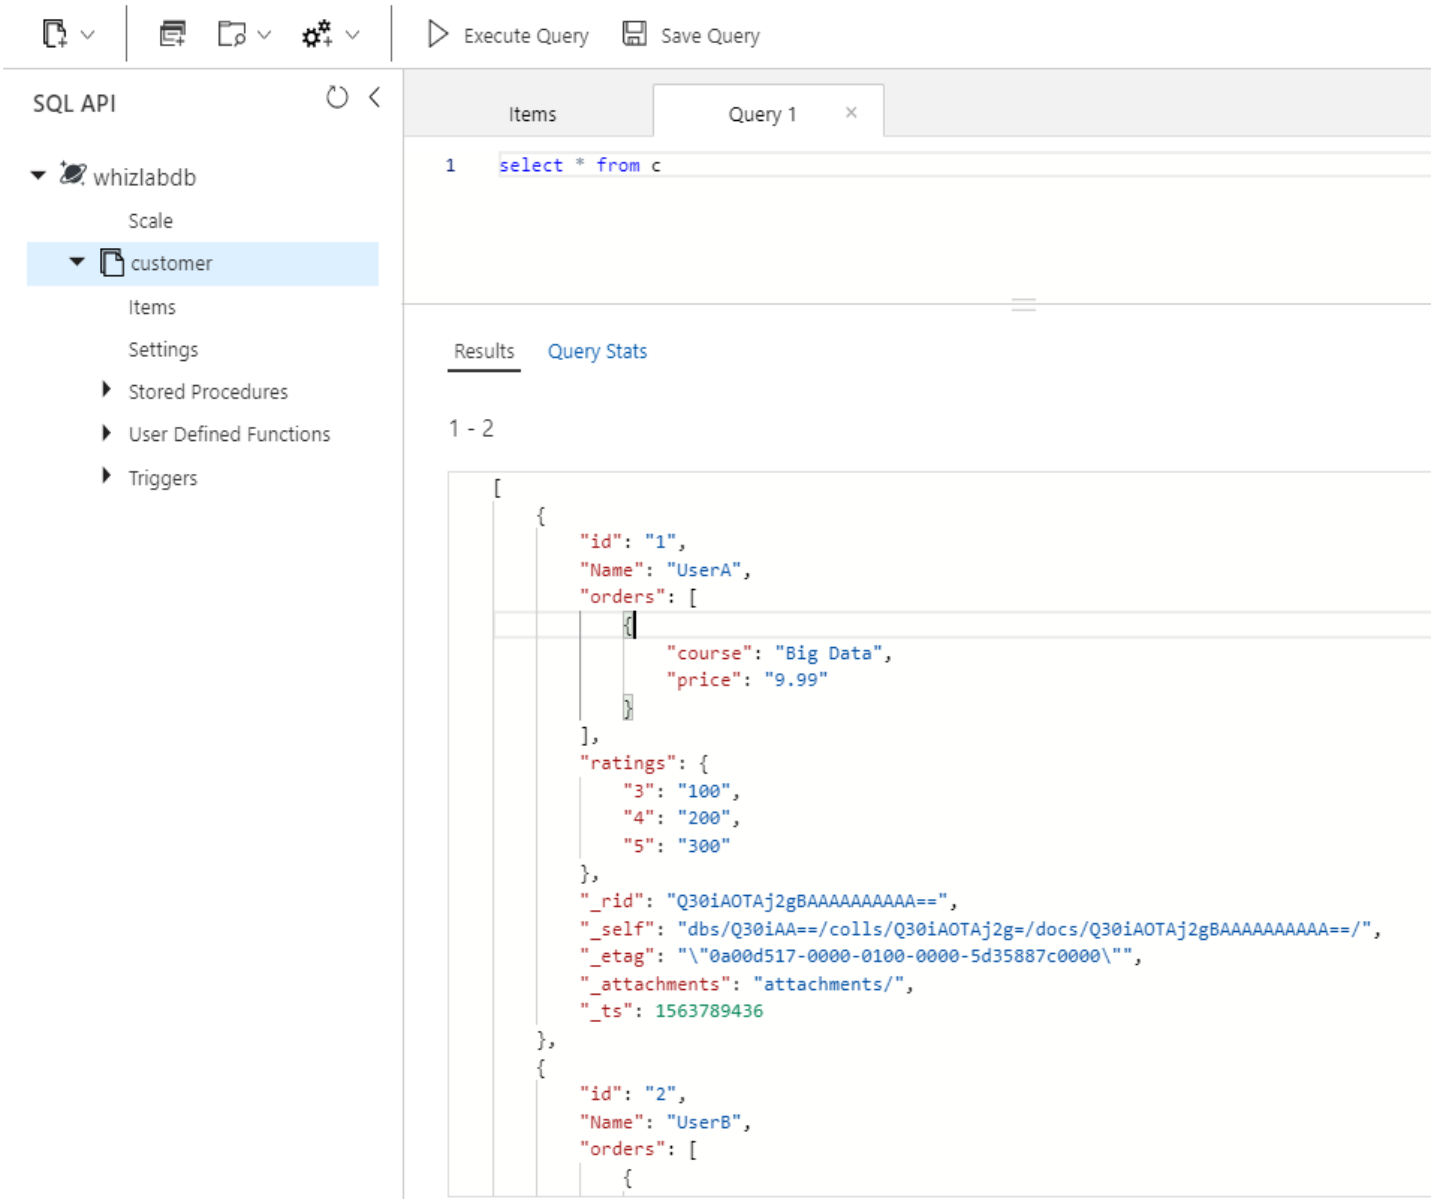

After executing the query

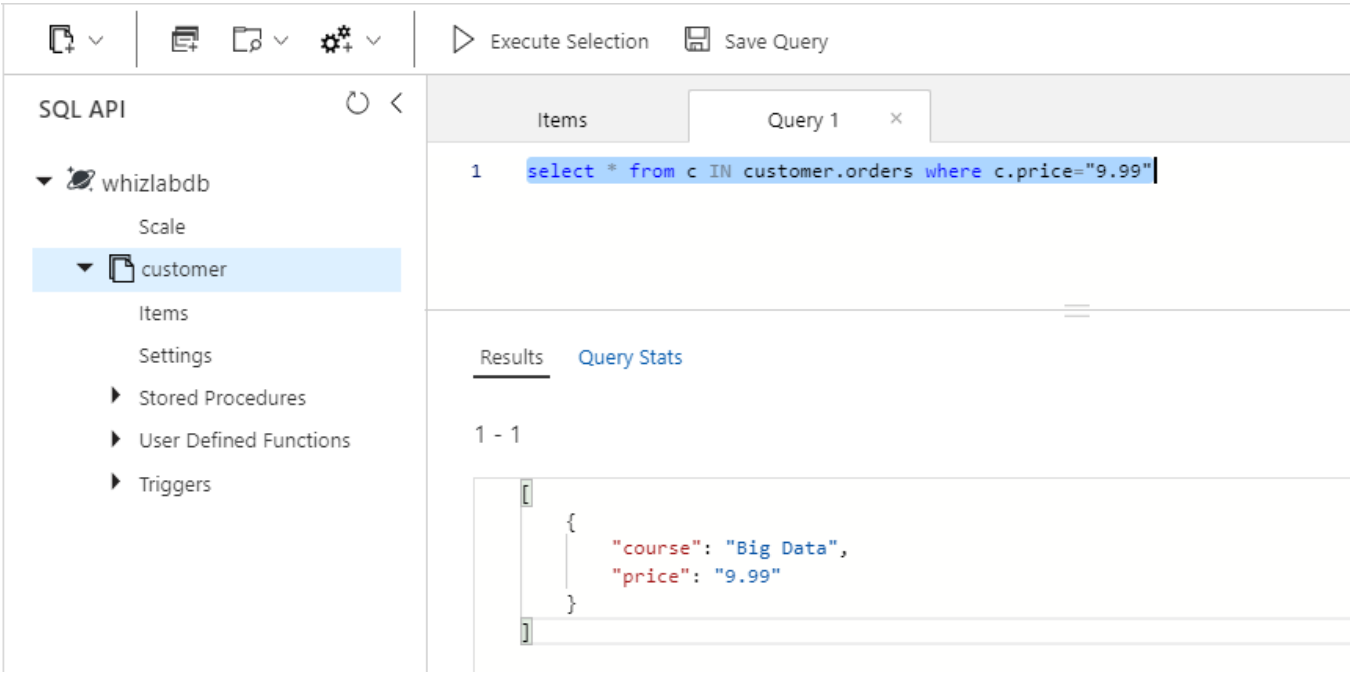

Since this is clear from the implementation, all other options are invalid

For more information on SQL queries for array objects, please visit the below URL

https://docs.microsoft.com/en-us/azure/cosmos-db/sql-query-object-array

Ask our Experts Rate this Question?

View Queries open

Question 32 Correct

Domain :Develop for Azure storage

View Case Study

You need to formulate a query that would be used to get all items where the order price is $9.99

```
select * from c IN [Slot1] where [Slot2] ="9.99"
```

Which of the following would go into Slot2?

A. orders

B. course

C. **c.price**

D. **c.orders**

Explanation:

Answer – C

Here since we need to filter based on the price, we should use the where clause to query based on the price

An example of an implementation of querying the data is given below

Data in the container

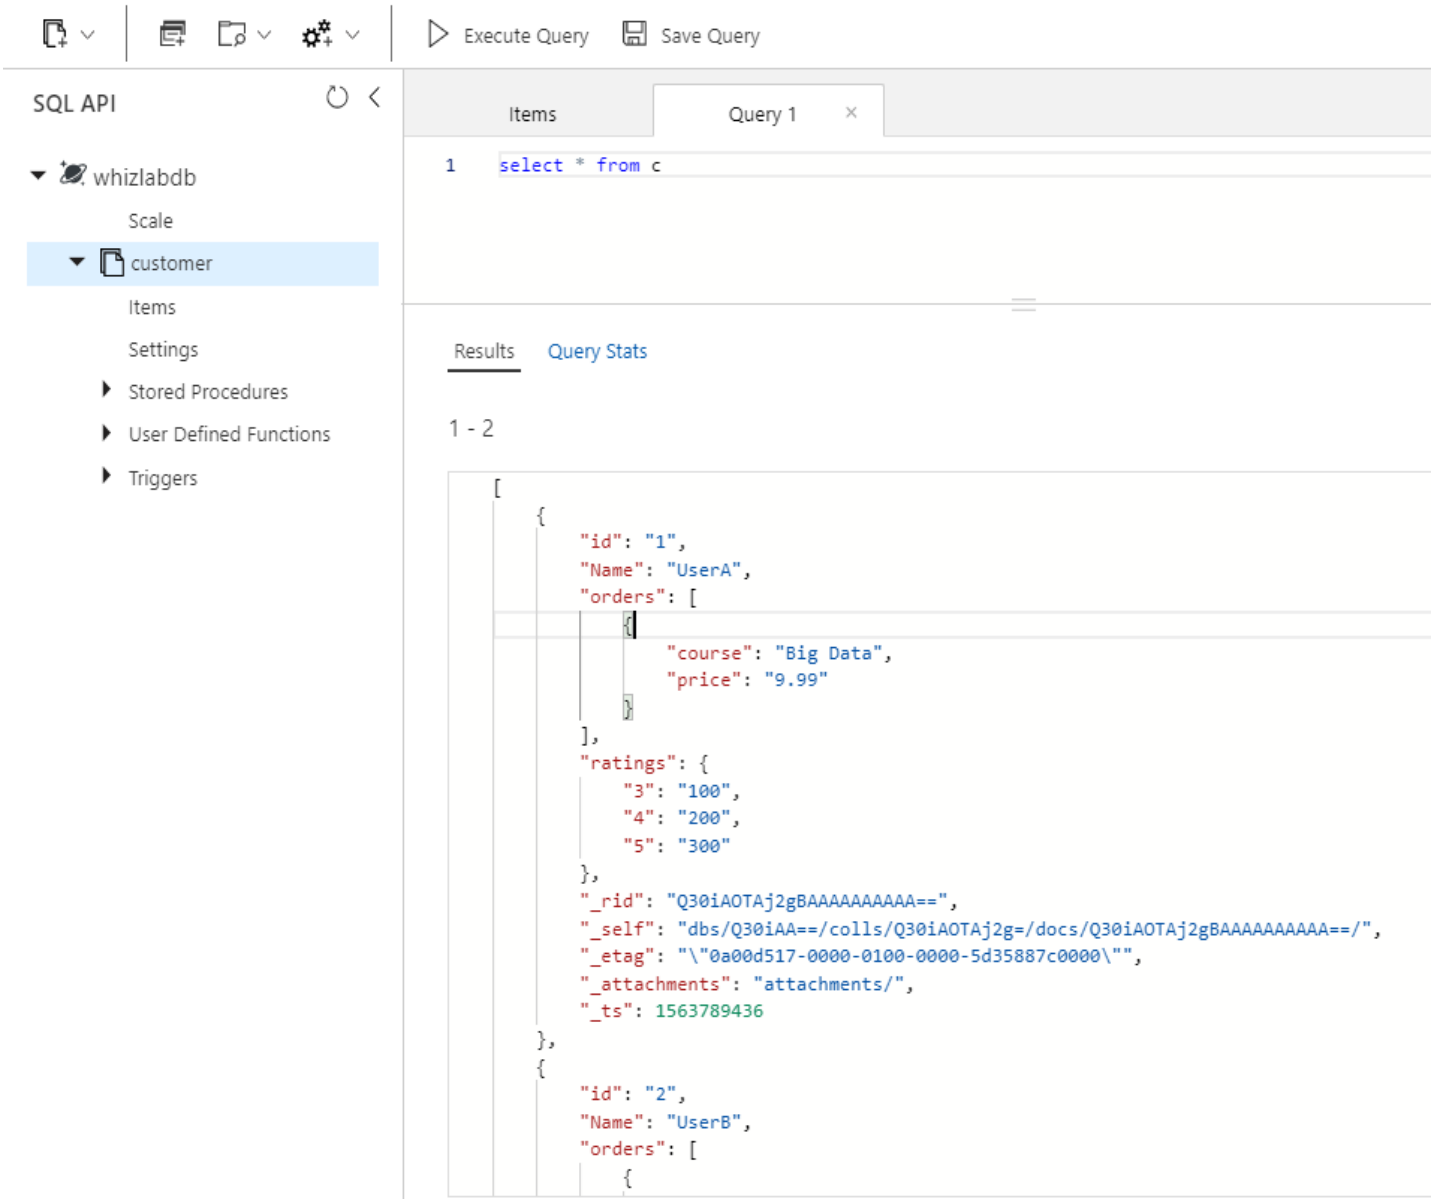

After executing the query

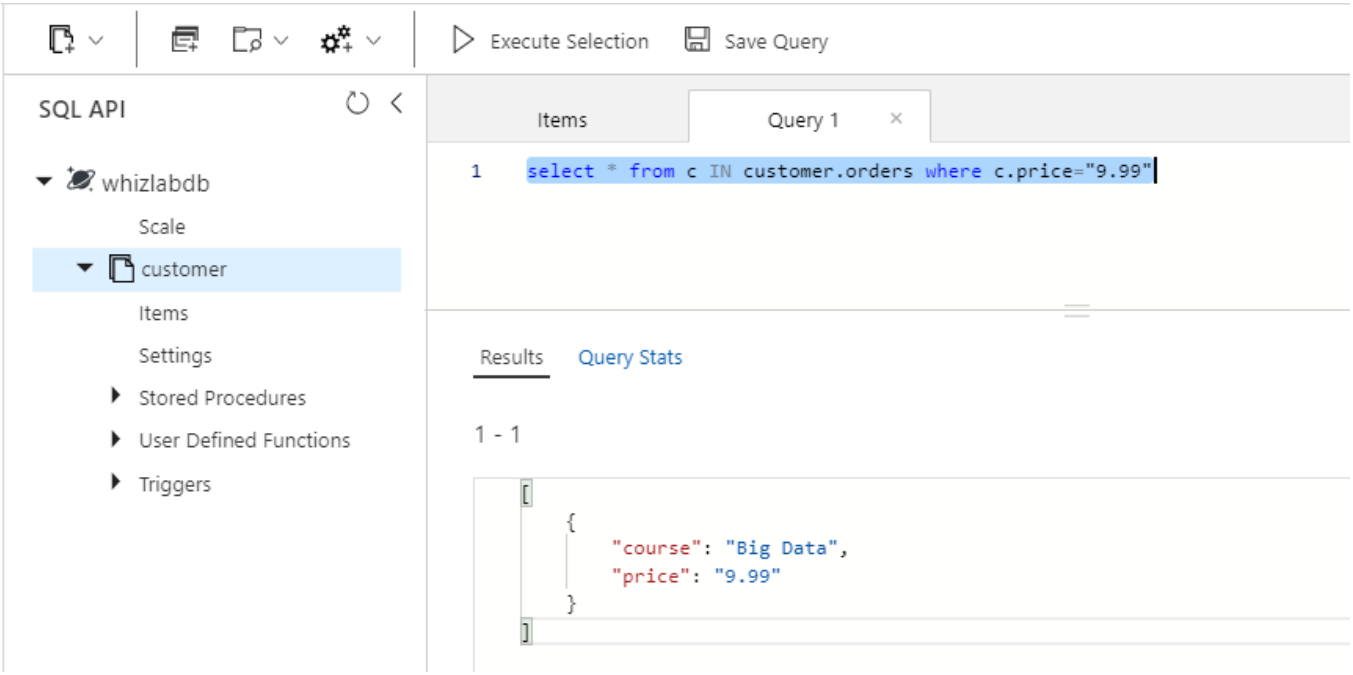

Since this is clear from the implementation, all other options are invalid

For more information on working with the WHERE clause, please visit the below URL

https://docs.microsoft.com/en-us/azure/cosmos-db/sql-query-where

Ask our Experts Rate this Question?

View Queries open

Question 33 Correct

Domain :Develop for Azure storage

View Case Study

You have to update the below code snippet that would be used to upload images to the Blob container.
Which of the following would go into Slot1?

```
Slot1 whizlabcloudBlockBlob =
cloudBlobContainer.GetBlockBlobReference(imgName);
await Slot2 .UploadFromFileAsync(imgFile);
```

- ✓ A. **CloudBlockBlob**
- B. **BlockBlob**
- C. **CloudBlob**
- D. **Blob**

Explanation:

Answer – A

The right data type is "CloudBlockBlob"

The Microsoft documentation gives an example on uploading content to Blob containers

Upload blobs to a container

The following code snippet gets a reference to a **CloudBlockBlob** object by calling the GetBlockBlobReference method on the container created in the previous section. It then uploads the selected local file to the blob by calling the UploadFromFileAsync method. This method creates the blob if it doesn't already exist, and overwrites it if it does.

C# Copy

```
// Create a file in your local MyDocuments folder to upload to a blob.
string localPath = Environment.GetFolderPath(Environment.SpecialFolder.MyDocuments);
string localFileName = "QuickStart_" + Guid.NewGuid().ToString() + ".txt";
string sourceFile = Path.Combine(localPath, localFileName);
// Write text to the file.
File.WriteAllText(sourceFile, "Hello, World!");

Console.WriteLine("Temp file = {0}", sourceFile);
Console.WriteLine("Uploading to Blob storage as blob '{0}'", localFileName);

// Get a reference to the blob address, then upload the file to the blob.
// Use the value of localFileName for the blob name.
CloudBlockBlob cloudBlockBlob = cloudBlobContainer.GetBlockBlobReference(localFileName);
await cloudBlockBlob.UploadFromFileAsync(sourceFile);
```

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on working with the Blob service from .Net, please visit the below URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-quickstart-blobs-dotnet

Ask our Experts Rate this Question? ☺ ☹

View Queries open

Question 34 Correct

Domain :Develop for Azure storage

View Case Study

You have to update the below code snippet that would be used to upload images to the Blob container.
Which of the following would go into Slot2?

```
Slot1 whizlabcloudBlockBlob =
cloudBlobContainer.GetBlockBlobReference(imgName);
await Slot2 .UploadFromFileAsync(imgFile);
```

- ✓ A. whizlabcloudBlockBlob
- B. BlockBlob
- C. CloudBlockBlob
- D. Blob

Explanation:

Answer – A

The right data type is "whizlabcloudBlockBlob"

The Microsoft documentation gives an example on uploading content to Blob containers

Upload blobs to a container

The following code snippet gets a reference to a **CloudBlockBlob** object by calling the GetBlockBlobReference method on the container created in the previous section. It then uploads the selected local file to the blob by calling the UploadFromFileAsync method. This method creates the blob if it doesn't already exist, and overwrites it if it does.

C# Copy

```
// Create a file in your local MyDocuments folder to upload to a blob.
string localPath = Environment.GetFolderPath(Environment.SpecialFolder.MyDocuments);
string localFileName = "QuickStart_" + Guid.NewGuid().ToString() + ".txt";
string sourceFile = Path.Combine(localPath, localFileName);
// Write text to the file.
File.WriteAllText(sourceFile, "Hello, World!");

Console.WriteLine("Temp file = {0}", sourceFile);
Console.WriteLine("Uploading to Blob storage as blob '{0}'", localFileName);

// Get a reference to the blob address, then upload the file to the blob.
// Use the value of localFileName for the blob name.
CloudBlockBlob cloudBlockBlob = cloudBlobContainer.GetBlockBlobReference(localFileName);
await cloudBlockBlob.UploadFromFileAsync(sourceFile);
```

Since this is clear from the Microsoft documentation, all other options are incorrect

For more information on working with the Blob service from .Net, please visit the below URL

https://docs.microsoft.com/en-us/azure/storage/blobs/storage-quickstart-blobs-dotnet

Ask our Experts Rate this Question?

View Queries open

Question 35 Correct

Domain :Develop Azure compute solutions

View Case Study

The Azure Web App needs to be running at all times. You have to ensure the most effective plan is assigned to the Azure Web App.
You decide to make the Azure Web App part of the “Shared” App Service Plan?
Would this fulfil the requirement?

A. Yes

✓ B. No

Explanation:

Answer – B

To fulfil this requirement, the App service plan must support the "Always On" feature. And this is not supported with the Shared App Service Plan

Below is the snippet of the features of the various App Service Plans

	FREE	SHARED	BASIC	STANDARD	PREMIUM	ISOLATED *	APP SERVICE LINUX	CONSUMPTION PLAN (FUNCTIONS)
– Settings								
64-bit			✓	✓	✓	✓	✓	✓
App Service Advisor *			✓	✓	✓	✓	✓	
Always On			✓	✓	✓	✓		

For more information on the App Service Plans, please visit the below URL

https://azure.microsoft.com/en-us/pricing/details/app-service/plans/

Ask our Experts Rate this Question?

View Queries open

Question 36 Incorrect

Domain :Develop Azure compute solutions

View Case Study

The Azure Web App needs to be running at all times. You have to ensure the most effective plan is assigned to the Azure Web App.
You decide to make the Azure Web App part of the "Basic" App Service Plan?
Would this fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – A

To fulfil this requirement, the App service plan must support the "Always On" feature. And this is supported with the Basic App Service Plan

Below is the snippet of the features of the various App Service Plans

	FREE	SHARED	BASIC	STANDARD	PREMIUM	ISOLATED *	APP SERVICE LINUX	CONSUMPTION PLAN (FUNCTIONS)
– Settings								
64-bit			✓	✓	✓	✓	✓	✓
App Service Advisor *			✓	✓	✓	✓	✓	
Always On			✓	✓	✓	✓		

For more information on the App Service Plans, please visit the below URL

https://azure.microsoft.com/en-us/pricing/details/app-service/plans/

Ask our Experts Rate this Question?

View Queries open

Question 37 Incorrect

Domain :Develop Azure compute solutions

View Case Study

The Azure Web App needs to be running at all times. You have to suggest the most cost effective plan for the Azure Web App. You recommended "Standard" as Web App Service Plan.
Would this fulfil the requirement?

A. Yes

B. No

Explanation:

Answer – B

To fulfil this requirement, the App service plan must support the "Always On" feature. And this is supported with the Standard App Service Plan. But this is not the most cost-effective App Service plan. Since the Basic App Service Plan already has this feature, you should choose that plan to cut on costs.

Below is the snippet of the features of the various App Service Plans

	FREE	SHARED	BASIC	STANDARD	PREMIUM	ISOLATED *	APP SERVICE LINUX	CONSUMPTION PLAN (FUNCTIONS)
– Settings								
64-bit			✓	✓	✓	✓	✓	✓
App Service Advisor *			✓	✓	✓	✓	✓	
Always On			✓	✓	✓	✓		

For more information on the App Service Plans, please visit the below URL

https://azure.microsoft.com/en-us/pricing/details/app-service/plans/

Ask our Experts Rate this Question?

View Queries open

Question 38 Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

View Case Study

A developer needs to enable Application Insights Profiler for the Azure Web App. Which of the following feature required to enable Application Insights Profiler for a Web App?

A. CORS configuration

✓ B. Always On setting

C. Enable Identity

D. Enable Custom domains

Explanation:

Answer - B

The requirement is to enable Application Insights Profiler for a web app which captures the data automatically at scale without negatively affecting to end user and widely used for monitoring purpose. and the given options are

A. CORS configuration

1. Cross-origin resource sharing (CORS) defines a way for client web applications that are loaded in one domain to interact with resources
2. Hence this is not going to help us to enable Application Insights Profiler

B. Always On setting

1. If a web app sit idle for too long, system unloads the website, and when traffic return, system need to load the Web App which causes longer response time and higher utilization of resources. By enabling 'Always On' setting (available for Standard tier websites), keeps the Web App up and running, which translates to higher availability and faster response times across the board.
2. Since requirement is to enable 'Application Insights Profiler' for a web App, Enabling 'Always On' is correct prerequisite.
3. Hence this is correct answer. We recommend ou to find detail at https://docs.microsoft.com/en-us/azure/azure-monitor/app/profiler

C. Enable Identity

1. identity is part of Azure Active Directory hence not related to Application Insights Profiler

D. Enable Custom domains

1. Enable Custom Domain is a feature to access any Azure URI by different or customized URI. Hence this is not related to enabling Application Insights Profiler

We have to enable the “Always On” setting.

This is also given in the Microsoft documentation

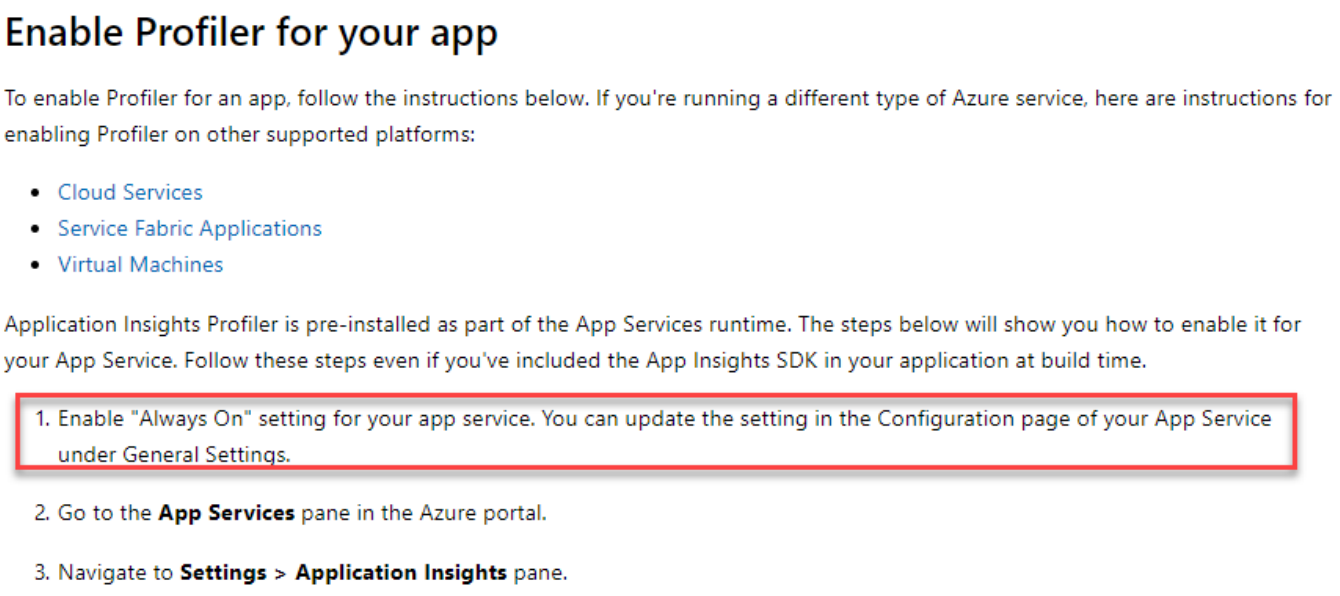

Enable Profiler for your app

To enable Profiler for an app, follow the instructions below. If you're running a different type of Azure service, here are instructions for enabling Profiler on other supported platforms:

- Cloud Services
- Service Fabric Applications
- Virtual Machines

Application Insights Profiler is pre-installed as part of the App Services runtime. The steps below will show you how to enable it for your App Service. Follow these steps even if you've included the App Insights SDK in your application at build time.

1. Enable "Always On" setting for your app service. You can update the setting in the Configuration page of your App Service under General Settings.
2. Go to the **App Services** pane in the Azure portal.
3. Navigate to **Settings** > **Application Insights** pane.

Since this is clearly given in the documentation, all other options are incorrect

For more information on enabling Application Insights for the Azure Web App, please visit the below URL

https://docs.microsoft.com/en-us/azure/azure-monitor/app/profiler

Ask our Experts Rate this Question?

View Queries open

Question 39 Correct

Domain :Develop for Azure storage

View Case Study

A lease needs to be applied on common blob's in the Azure storage account.
How would you complete the below REST API call for implementing a blob lease?
https://whizlabstore2020.blob.core.windows.net/demo/whizlabcommon.json?

Slot1

Which of the following should go into Slot1?

A. getlease

✓ B. comp=lease

C. get=lease

D. obj=lease

Explanation:

Answer - B

We have to use the query string parameter as “comp=lease”

This is also given in the Microsoft documentation

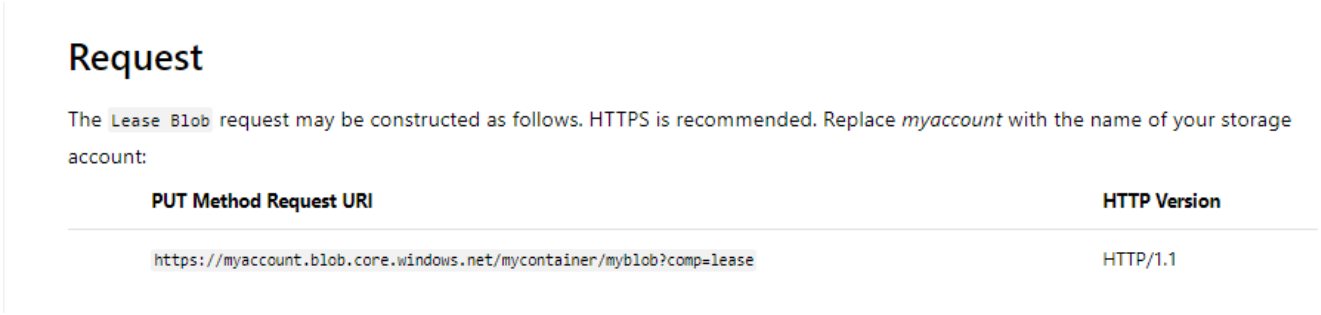

Request

The `Lease Blob` request may be constructed as follows. HTTPS is recommended. Replace *myaccount* with the name of your storage account:

PUT Method Request URI	HTTP Version
`https://myaccount.blob.core.windows.net/mycontainer/myblob?comp=lease`	HTTP/1.1

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with Blob lease, please visit the below URL

https://docs.microsoft.com/en-us/rest/api/storageservices/lease-blob

Ask our Experts Rate this Question?

View Queries open

Question 40 Correct

Domain :Develop for Azure storage

A company is developing a series of applications. Each of these applications would be interacting with separate Azure CosmosDB accounts. Each application has a different requirement when it comes to accessing the underlying data. You have to set the consistency level for the Azure CosmosDB accounts based on each application requirement. You have to choose the most cost-effective consistency level for each CosmosDB account. Below are the requirements for each application when it comes to the consistency of the underlying data

Application Name	Data requirement
whizlabappA	This is not a critical application hence here no ordering

	of reads is required.
whizlabappB	Here the application uses must never see out-of-order writes
whizlabappC	Here the users must always see the latest committed write.
whizlabappD	Here the data can be stale by at most 2 versions

Which of the following would you choose as the consistency level for the CosmosDB account used by the application "whizlabappA"?

A. **Strong**

B. **Bounded Staleness**

C. **Session**

D. **Consistent prefix**

✓ E. **Eventual**

Explanation:

Answer – E

Since the requirement is that the users don't mind seeing out of order reads, one can use the Eventual consistency level as the most cost-effective consistency level for the CosmosDB account.

The Microsoft documentation mentions the following

Guarantees associated with consistency levels

The comprehensive SLAs provided by Azure Cosmos DB guarantee that 100 percent of read requests meet the consistency guarantee for any consistency level you choose. A read request meets the consistency SLA if all the consistency guarantees associated with the consistency level are satisfied. The precise definitions of the five consistency levels in Azure Cosmos DB using the TLA+ specification language are provided in the azure-cosmos-tla GitHub repo.

The semantics of the five consistency levels are described here:

- **Strong**: Strong consistency offers a linearizability guarantee. Linearizability refers to serving requests concurrently. The reads are guaranteed to return the most recent committed version of an item. A client never sees an uncommitted or partial write. Users are always guaranteed to read the latest committed write.
- **Bounded staleness**: The reads are guaranteed to honor the consistent-prefix guarantee. The reads might lag behind writes by at most "*K*" versions (i.e., "updates") of an item or by "*T*" time interval. In other words, when you choose bounded staleness, the "staleness" can be configured in two ways:
  - The number of versions (*K*) of the item
  - The time interval (*T*) by which the reads might lag behind the writes

  Bounded staleness offers total global order except within the "staleness window." The monotonic read guarantees exist within a region both inside and outside the staleness window. Strong consistency has the same semantics as the one offered by bounded staleness. The staleness window is equal to zero. Bounded staleness is also referred to as time-delayed linearizability. When a client performs read operations within a region that accepts writes, the guarantees provided by bounded staleness consistency are identical to those guarantees by the strong consistency.
- **Session**: Within a single client session reads are guaranteed to honor the consistent-prefix (assuming a single "writer" session), monotonic reads, monotonic writes, read-your-writes, and write-follows-reads guarantees. Clients outside of the session performing writes will see eventual consistency.
- **Consistent prefix**: Updates that are returned contain some prefix of all the updates, with no gaps. Consistent prefix consistency level guarantees that reads never see out-of-order writes.
- **Eventual**: There's no ordering guarantee for reads. In the absence of any further writes, the replicas eventually converge.

Since this is clearly given in the documentation, all other options are incorrect

For more information on CosmosDB consistency levels, one can go to the below URL

https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels

Ask our Experts Rate this Question?

View Queries open

Question 41 Incorrect

Domain :Develop for Azure storage

A company is developing a series of applications. Each of these applications would be interacting with separate Azure CosmosDB accounts. Each application has a different requirement when it comes to accessing the underlying data. You have to set the consistency level for the Azure CosmosDB accounts based on each application requirement. You have to choose the most cost-effective consistency level for each CosmosDB account. Below are the requirements for each application when it comes to the consistency of the underlying data

Application Name	Data requirement
whizlabappA	This is not a critical application hence here no ordering of reads is required.
whizlabappB	Here the application uses must never see out-of-order writes
whizlabappC	Here the users must always see the latest committed write.
whizlabappD	Here the data can be stale by at most 2 versions

Which of the following would you choose as the consistency level for the CosmosDB account used by the application "whizlabappB"?

- A. **Strong** ✗
- B. **Bounded Staleness**
- C. **Session**
- D. **Consistent prefix** ✓
- E. **Eventual**

Explanation:

Answer - D

Since here the requirement is that the user must never see out-of-order writes, the most cost-effective option is to use the "Consistent prefix" consistency level.

The Microsoft documentation mentions the following

Guarantees associated with consistency levels

The comprehensive SLAs provided by Azure Cosmos DB guarantee that 100 percent of read requests meet the consistency guarantee for any consistency level you choose. A read request meets the consistency SLA if all the consistency guarantees associated with the consistency level are satisfied. The precise definitions of the five consistency levels in Azure Cosmos DB using the TLA+ specification language are provided in the azure-cosmos-tla GitHub repo.

The semantics of the five consistency levels are described here:

- **Strong**: Strong consistency offers a linearizability guarantee. Linearizability refers to serving requests concurrently. The reads are guaranteed to return the most recent committed version of an item. A client never sees an uncommitted or partial write. Users are always guaranteed to read the latest committed write.
- **Bounded staleness**: The reads are guaranteed to honor the consistent-prefix guarantee. The reads might lag behind writes by at most "*K*" versions (i.e., "updates") of an item or by "*T*" time interval. In other words, when you choose bounded staleness, the "staleness" can be configured in two ways:
  - The number of versions (*K*) of the item
  - The time interval (*T*) by which the reads might lag behind the writes

  Bounded staleness offers total global order except within the "staleness window." The monotonic read guarantees exist within a region both inside and outside the staleness window. Strong consistency has the same semantics as the one offered by bounded staleness. The staleness window is equal to zero. Bounded staleness is also referred to as time-delayed linearizability. When a client performs read operations within a region that accepts writes, the guarantees provided by bounded staleness consistency are identical to those guarantees by the strong consistency.
- **Session**: Within a single client session reads are guaranteed to honor the consistent-prefix (assuming a single "writer" session), monotonic reads, monotonic writes, read-your-writes, and write-follows-reads guarantees. Clients outside of the session performing writes will see eventual consistency.
- **Consistent prefix**: Updates that are returned contain some prefix of all the updates, with no gaps. Consistent prefix consistency level guarantees that reads never see out-of-order writes.
- **Eventual**: There's no ordering guarantee for reads. In the absence of any further writes, the replicas eventually converge.

Since this is clearly given in the documentation, all other options are incorrect

For more information on CosmosDB consistency levels, one can go to the below URL

https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels

Ask our Experts

Rate this Question?

View Queries open

Question 42 Correct

Domain :Develop for Azure storage

A company is developing a series of applications. Each of these applications would be interacting with separate Azure CosmosDB accounts. Each application has a different requirement when it comes to accessing the underlying data. You have to set the consistency level for the Azure CosmosDB accounts based on each application requirement. You have to choose the most cost-effective consistency level for each CosmosDB account. Below are the requirements for each application when it comes to the consistency of the underlying data

Application Name	Data requirement
whizlabappA	This is not a critical application hence here no ordering of reads is required.
whizlabappB	Here the application uses must never see out-of-order writes
whizlabappC	Here the users must always see the latest committed write.
whizlabappD	Here the data can be stale by at most 2 versions

Which of the following would you choose as the consistency level for the CosmosDB account used by the application "whizlabappC"?

- A. **Strong**
- B. **Bounded Staleness**
- C. **Session**
- D. **Consistent prefix**
- E. **Eventual**

Explanation:

Answer – A

Since here we need to ensure that the user always sees the latest committed version, we have to choose the "Strong" consistency level.

The Microsoft documentation mentions the following

Guarantees associated with consistency levels

The comprehensive SLAs provided by Azure Cosmos DB guarantee that 100 percent of read requests meet the consistency guarantee for any consistency level you choose. A read request meets the consistency SLA if all the consistency guarantees associated with the consistency level are satisfied. The precise definitions of the five consistency levels in Azure Cosmos DB using the TLA+ specification language are provided in the azure-cosmos-tla GitHub repo.

The semantics of the five consistency levels are described here:

- **Strong**: Strong consistency offers a linearizability guarantee. Linearizability refers to serving requests concurrently. The reads are guaranteed to return the most recent committed version of an item. A client never sees an uncommitted or partial write. Users are always guaranteed to read the latest committed write.
- **Bounded staleness**: The reads are guaranteed to honor the consistent-prefix guarantee. The reads might lag behind writes by at most "*K*" versions (i.e., "updates") of an item or by "*T*" time interval. In other words, when you choose bounded staleness, the "staleness" can be configured in two ways:
  - The number of versions (*K*) of the item
  - The time interval (*T*) by which the reads might lag behind the writes

  Bounded staleness offers total global order except within the "staleness window." The monotonic read guarantees exist within a region both inside and outside the staleness window. Strong consistency has the same semantics as the one offered by bounded staleness. The staleness window is equal to zero. Bounded staleness is also referred to as time-delayed linearizability. When a client performs read operations within a region that accepts writes, the guarantees provided by bounded staleness consistency are identical to those guarantees by the strong consistency.
- **Session**: Within a single client session reads are guaranteed to honor the consistent-prefix (assuming a single "writer" session), monotonic reads, monotonic writes, read-your-writes, and write-follows-reads guarantees. Clients outside of the session performing writes will see eventual consistency.
- **Consistent prefix**: Updates that are returned contain some prefix of all the updates, with no gaps. Consistent prefix consistency level guarantees that reads never see out-of-order writes.
- **Eventual**: There's no ordering guarantee for reads. In the absence of any further writes, the replicas eventually converge.

Since this is clearly given in the documentation, all other options are incorrect

For more information on CosmosDB consistency levels, one can go to the below URL

https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels

Ask our Experts

Rate this Question?

View Queries open

Question 43 Correct

Domain :Develop for Azure storage

A company is developing a series of applications. Each of these applications would be interacting with separate Azure CosmosDB accounts. Each application has a different requirement when it comes to accessing the underlying data. You have to set the consistency level for the Azure CosmosDB accounts based on each application requirement. You have to choose the most cost-effective consistency level for each CosmosDB account. Below are the requirements for each application when it comes to the consistency of the underlying data

Application Name	Data requirement
whizlabappA	This is not a critical application hence here no ordering of reads is required.
whizlabappB	Here the application uses must never see out-of-order writes
whizlabappC	Here the users must always see the latest committed write.
whizlabappD	Here the data can be stale by at most 2 versions

Which of the following would you choose as the consistency level for the CosmosDB account used by the application "whizlabappD"?

A. **Strong**

✓ B. **Bounded Staleness**

C. **Session**

D. **Consistent prefix**

E. **Eventual**

Explanation:

Answer – B

Here since we have a staleness of data by a set version count, we can use the Bounded staleness consistency level.

The Microsoft documentation mentions the following

Guarantees associated with consistency levels

The comprehensive SLAs provided by Azure Cosmos DB guarantee that 100 percent of read requests meet the consistency guarantee for any consistency level you choose. A read request meets the consistency SLA if all the consistency guarantees associated with the consistency level are satisfied. The precise definitions of the five consistency levels in Azure Cosmos DB using the TLA+ specification language are provided in the azure-cosmos-tla GitHub repo.

The semantics of the five consistency levels are described here:

- **Strong**: Strong consistency offers a linearizability guarantee. Linearizability refers to serving requests concurrently. The reads are guaranteed to return the most recent committed version of an item. A client never sees an uncommitted or partial write. Users are always guaranteed to read the latest committed write.
- **Bounded staleness**: The reads are guaranteed to honor the consistent-prefix guarantee. The reads might lag behind writes by at most "*K*" versions (i.e., "updates") of an item or by "*T*" time interval. In other words, when you choose bounded staleness, the "staleness" can be configured in two ways:
  - The number of versions (*K*) of the item
  - The time interval (*T*) by which the reads might lag behind the writes

  Bounded staleness offers total global order except within the "staleness window." The monotonic read guarantees exist within a region both inside and outside the staleness window. Strong consistency has the same semantics as the one offered by bounded staleness. The staleness window is equal to zero. Bounded staleness is also referred to as time-delayed linearizability. When a client performs read operations within a region that accepts writes, the guarantees provided by bounded staleness consistency are identical to those guarantees by the strong consistency.
- **Session**: Within a single client session reads are guaranteed to honor the consistent-prefix (assuming a single "writer" session), monotonic reads, monotonic writes, read-your-writes, and write-follows-reads guarantees. Clients outside of the session performing writes will see eventual consistency.
- **Consistent prefix**: Updates that are returned contain some prefix of all the updates, with no gaps. Consistent prefix consistency level guarantees that reads never see out-of-order writes.
- **Eventual**: There's no ordering guarantee for reads. In the absence of any further writes, the replicas eventually converge.

Since this is clearly given in the documentation, all other options are incorrect

For more information on CosmosDB consistency levels, one can go to the below URL

https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels

Ask our Experts

Rate this Question?

View Queries open

Question 44 Correct

Domain :Develop Azure compute solutions

A company has a web application deployed to Azure. The web application is currently being hosted as part of the Azure Web App service. There is a requirement to stream the logs from the web app and filter out on any errors.
You have to complete the below Azure CLI command for this requirement

az webapp Slot1 Slot2 --name “whizlabapp”

--resource-group “whizlab-rg” Slot3 Error

Which of the following would go into Slot1?

A. file

✓ B. log

C. tail

D. stream

Explanation:

Answer – B

Since we need to get the log files for the Azure Web app, we have to use the “log” option in the command.

An example of this is also given in the Microsoft documentation

Streaming with Azure CLI

To stream logging information, open a new command prompt, PowerShell, Bash, or Terminal session and enter the following command:

Copy

```
az webapp log tail --name appname --resource-group myResourceGroup
```

This command connects to the app named 'appname' and begin streaming information to the window as log events occur on the app. Any information written to files ending in .txt, .log, or .htm that are stored in the /LogFiles directory (d:/home/logfiles) is streamed to the local console.

To filter specific events, such as errors, use the **--Filter** parameter. For example:

Copy

```
az webapp log tail --name appname --resource-group myResourceGroup --filter Error
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with diagnostics logs, one can go to the below URL

https://docs.microsoft.com/en-us/azure/app-service/troubleshoot-diagnostic-logs

Ask our Experts Rate this Question?

View Queries open

Question 45 Correct

Domain :Develop Azure compute solutions

A company has a web application deployed to Azure. The web application is currently being hosted as part of the Azure Web App service. There is a requirement to stream the logs from the web app and filter out on any errors.
You have to complete the below Azure CLI command for this requirement

az webapp Slot1 Slot2 --name “whizlabapp”

--resource-group “whizlab-rg” Slot3 Error

Which of the following would go into Slot2?

A. file

B. log

✓ C. tail

D. stream

Explanation:

Answer – C

Since we need to stream the log files , the next option to include is the “tail” option

An example of this is also given in the Microsoft documentation

Streaming with Azure CLI

To stream logging information, open a new command prompt, PowerShell, Bash, or Terminal session and enter the following command:

Copy

```
az webapp log tail --name appname --resource-group myResourceGroup
```

This command connects to the app named 'appname' and begin streaming information to the window as log events occur on the app. Any information written to files ending in .txt, .log, or .htm that are stored in the /LogFiles directory (d:/home/logfiles) is streamed to the local console.

To filter specific events, such as errors, use the **--Filter** parameter. For example:

Copy

```
az webapp log tail --name appname --resource-group myResourceGroup --filter Error
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with diagnostics logs, one can go to the below URL

https://docs.microsoft.com/en-us/azure/app-service/troubleshoot-diagnostic-logs

Ask our Experts Rate this Question?

View Queries open

Question 46 Correct

Domain :Develop Azure compute solutions

A company has a web application deployed to Azure. The web application is currently being hosted as part of the Azure Web App service. There is a requirement to stream the logs from the web app and filter out on any errors.
You have to complete the below Azure CLI command for this requirement

az webapp [Slot1] [Slot2] --name “whizlabapp”

--resource-group “whizlab-rg” [Slot3] Error

Which of the following would go into Slot3?

A. **--path**

B. **--filter**

C. **--key**

D. **--type**

Explanation:

Answer – B

Since we have the filter out on the errors, we can use the filter option for the command.

An example of this is also given in the Microsoft documentation

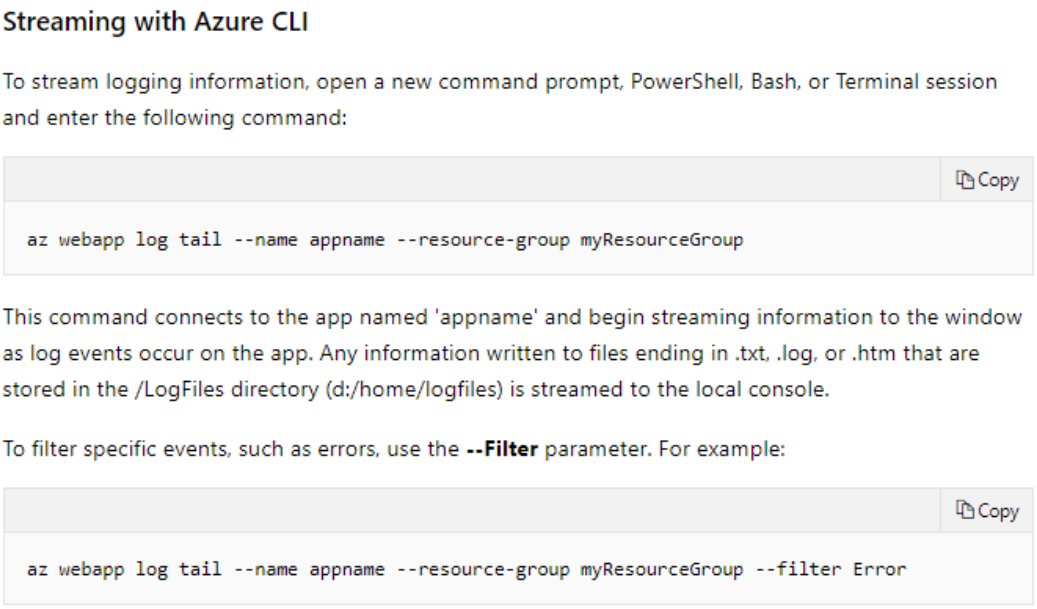

Streaming with Azure CLI

To stream logging information, open a new command prompt, PowerShell, Bash, or Terminal session and enter the following command:

Copy

```
az webapp log tail --name appname --resource-group myResourceGroup
```

This command connects to the app named 'appname' and begin streaming information to the window as log events occur on the app. Any information written to files ending in .txt, .log, or .htm that are stored in the /LogFiles directory (d:/home/logfiles) is streamed to the local console.

To filter specific events, such as errors, use the **--Filter** parameter. For example:

Copy

```
az webapp log tail --name appname --resource-group myResourceGroup --filter Error
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with diagnostics logs, one can go to the below URL

https://docs.microsoft.com/en-us/azure/app-service/troubleshoot-diagnostic-logs

Ask our Experts

Rate this Question?

View Queries

open

Question 47 Incorrect

Domain :Connect to and consume Azure services and third-party services

A team has to integrate various modules of an application with the Azure Event Grid service. They have to filter events which are sent to the various application endpoints. The requirements for the type of messages that need to be received by the different endpoints are given below

Application Endpoint	Message requirement
EndpointA	Receives failure messages for any resources deployed to the Azure subscription
EndpointB	Receives messages whenever objects are added to a specific container in Azure Blob storage
EndpointC	Receive messages whenever data fields in the message has the value of "Organization"

Which of the following would you use as a filter option for messages that need to be sent to EndpointA?

A. Subject begins with or ends with

B. Advanced fields and operators

✓ C. ResourceTypes ✗

D. EventTypes ✓

Explanation:

Answer - D

Since here we just need to filter on the event types itself, we can use the "EventTypes" filter.

The Microsoft documentation mentions the following

Event type filtering

By default, all event types for the event source are sent to the endpoint. You can decide to send only certain event types to your endpoint. For example, you can get notified of updates to your resources, but not notified for other operations like deletions. In that case, filter by the `Microsoft.Resources.ResourceWriteSuccess` event type. Provide an array with the event types, or specify `All` to get all event types for the event source.

The JSON syntax for filtering by event type is:

JSON Copy

```
"filter": {
  "includedEventTypes": [
    "Microsoft.Resources.ResourceWriteFailure",
    "Microsoft.Resources.ResourceWriteSuccess"
  ]
}
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on event filtering in Azure Event Grid, one can go to the below URL

https://docs.microsoft.com/en-us/azure/event-grid/event-filtering

Ask our Experts Rate this Question?

View Queries open

Question 48 Correct

Domain :Connect to and consume Azure services and third-party services

A team has to integrate various modules of an application with the Azure Event Grid service. They have to filter events which are sent to the various application endpoints. The requirements for the type of messages that need to be received by the different endpoints are given below

Application Endpoint	Message requirement
EndpointA	Receives failure messages for any resources deployed to the Azure subscription
EndpointB	Receives messages whenever objects are added to a specific container in Azure Blob storage
EndpointC	Receive messages whenever data fields in the message has the value of “Organization”

Which of the following would you use as a filter option for messages that need to be sent to EndpointB?

- ✓ A. **Subject begins with or ends with**
- B. **Advanced fields and operators**
- C. **ResourceTypes**
- D. **EventTypes**

Explanation:

Answer – A

Since here we need to check on the messages sent to a container, so we have to check the subject of the message

The Microsoft documentation mentions the following

Subject filtering

For simple filtering by subject, specify a starting or ending value for the subject. For example, you can specify the subject ends with `.txt` to only get events related to uploading a text file to storage account. Or, you can filter the subject begins with `/blobServices/default/containers/testcontainer` to get all events for that container but not other containers in the storage account.

When publishing events to custom topics, create subjects for your events that make it easy for subscribers to know whether they're interested in the event. Subscribers use the subject property to filter and route events. Consider adding the path for where the event happened, so subscribers can filter by segments of that path. The path enables subscribers to narrowly or broadly filter events. If you provide a three segment path like `/A/B/C` in the subject, subscribers can filter by the first segment `/A` to get a broad set of events. Those subscribers get events with subjects like `/A/B/C` or `/A/D/E`. Other subscribers can filter by `/A/B` to get a narrower set of events.

The JSON syntax for filtering by subject is:

JSON Copy

```
"filter": {
  "subjectBeginsWith": "/blobServices/default/containers/mycontainer/log",
  "subjectEndsWith": ".jpg"
}
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on event filtering in Azure Event Grid, one can go to the below URL

https://docs.microsoft.com/en-us/azure/event-grid/event-filtering

Ask our Experts Rate this Question?

View Queries open

Question 49 Correct

Domain :Connect to and consume Azure services and third-party services

A team has to integrate various modules of an application with the Azure Event Grid service. They have to filter events which are sent to the various application endpoints. The requirements for the type of messages that need to be received by the different endpoints are given below

Application Endpoint	Message requirement
EndpointA	Receives failure messages for any resources deployed to the Azure subscription
EndpointB	Receives messages whenever objects are added to a specific container in Azure Blob storage
EndpointC	Receive messages whenever data fields in the message has the value of "Organization"

Which of the following would you use as a filter option for messages that need to be sent to EndpointC?

A. **Subject begins with or ends with**

✓ B. **Advanced fields and operators**

C. **ResourceTypes**

D. **EventTypes**

Explanation:

Answer – B

Here since we need a more advanced scenario and check for the data field values, we have to choose the "Advanced fields and operators" filter option

The Microsoft documentation mentions the following

Advanced filtering

To filter by values in the data fields and specify the comparison operator, use the advanced filtering option. In advanced filtering, you specify the:

- operator type - The type of comparison.
- key - The field in the event data that you're using for filtering. It can be a number, boolean, or string.
- value or values - The value or values to compare to the key.

The JSON syntax for using advanced filters is:

JSON Copy

```
"filter": {
  "advancedFilters": [
    {
      "operatorType": "NumberGreaterThanOrEquals",
      "key": "Data.Key1",
      "value": 5
    },
    {
      "operatorType": "StringContains",
      "key": "Subject",
      "values": ["container1", "container2"]
    }
  ]
}
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on event filtering in Azure Event Grid, one can go to the below URL

https://docs.microsoft.com/en-us/azure/event-grid/event-filtering

Ask our Experts Rate this Question?

View Queries open

Question 50 Correct

Domain :Connect to and consume Azure services and third-party services

A development team have deployed an API management instance. An application sits behind the API management instance. The application accepts all data in JSON format. An external consultant currently connects to the API management instance. The data sent by the external consultant is in XML format.
You have to ensure the data gets converted to JSON by the API management instance.
You decide to implement an API management policy.
Would this fulfil the requirement?

✓ A. Yes

B. No

Explanation:

Answer – A

You can use a policy to convert the data.

The Microsoft documentation mentions the following

Convert XML to JSON

The `xml-to-json` policy converts a request or response body from XML to JSON. This policy can be used to modernize APIs based on XML-only backend web services.

Policy statement

XML Copy

```
<xml-to-json kind="javascript-friendly | direct" apply="always | content-type-xml" consider-a
```

Example

XML Copy

```
<policies>
    <inbound>
        <base />
    </inbound>
    <outbound>
        <base />
        <xml-to-json kind="direct" apply="always" consider-accept-header="false" />
    </outbound>
</policies>
```

For more information on API management policies, one can go to the below URL

https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies

Ask our Experts Rate this Question?

View Queries open

Question 51 Correct

Domain :Connect to and consume Azure services and third-party services

A development team have deployed an API management instance. An application sits behind the API management instance. The application accepts all data in JSON format. An external consultant currently connects to the API management instance. The data sent by the external consultant is in XML format.
You have to ensure the data gets converted to JSON by the API management instance.
You decide to create an Azure Event Hub namespace.
Would this fulfil the requirement?

A. Yes

✓ B. No

Explanation:

Answer – B

The Azure Event Hub could be used to log events from the Azure API Management instance

For more information on logging to Azure Event Hubs, one can go to the below URL

https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-log-event-hubs

Ask our Experts Rate this Question?

View Queries open

Question 52 Correct

Domain :Connect to and consume Azure services and third-party services

A development team have deployed an API management instance. An application sits behind the API management instance. The application accepts all data in JSON format. An external consultant currently connects to the API management instance. The data sent by the external consultant is in XML format.
You have to ensure the data gets converted to JSON by the API management instance.
You decide to implement RBAC.
Would this fulfil the requirement?

A. Yes

✓ B. No

Explanation:

Answer – B

RBAC is used to provide Role Based Access Control in Azure API Management

For more information on role-based access control for Azure API management, one can go to the below URL

https://docs.microsoft.com/en-us/azure/api-management/api-management-role-based-access-control

Ask our Experts Rate this Question?

View Queries open

Question 53 Incorrect

Domain :Connect to and consume Azure services and third-party services

A development team needs to develop an application module that needs to interact with an Azure service bus queue. Below is the snippet of code that needs to be completed. The code is used to send a message to the queue.

```
        const string whizlabServiceBusConnectionString = "...";
        const string QueueName = "whizlabqueue";
        static IQueueClient whizlabqueueClient;
        static async Task SideSync()
        {
            whizlabqueueClient = new    Slot1    (ServiceBusConnectionString,
QueueName);
            await SendMessagesAsync(numberOfMessages);
            await queueClient.CloseAsync();
        }

        static async Task SendMessagesAsync(string messageBody)
        {
            try
            {
                    var message = new
Message(Encoding.UTF8.GetBytes(messageBody));
                    await queueClient.    Slot2    (message);
                }
            }
            catch (Exception whizlabexception)
            {
                Console.WriteLine($"{DateTime.Now} :: Exception:
{whizlabexception.Message}");
            } } } }
```

Which of the following needs to go into Slot1?

A. Client

B. ServiceBusClient

C. QueueClient

D. BusClient

Explanation:

Answer – C

Since we need to interact with a queue, we have to use the "QueueClient" class.

An example of this is given in the Microsoft documentation

4. Directly after `Main()`, add the following asynchronous `MainAsync()` method that calls the send messages method:

C# Copy

```
static async Task MainAsync()
{
    const int numberOfMessages = 10;
    queueClient = new QueueClient(ServiceBusConnectionString, QueueName);

    Console.WriteLine("======================================================");
    Console.WriteLine("Press ENTER key to exit after sending all the messages.");
    Console.WriteLine("======================================================");

    // Send messages.
    await SendMessagesAsync(numberOfMessages);

    Console.ReadKey();

    await queueClient.CloseAsync();
}
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with queues in .Net, one can go to the below URL

https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-dotnet-get-started-with-queues

Ask our Experts Rate this Question?

View Queries open

Question 54 Correct

Domain :Connect to and consume Azure services and third-party services

A development team needs to develop an application module that needs to interact with an Azure service bus queue. Below is the snippet of code that needs to be completed. The code is used to send a message to the queue.

```
        const string whizlabServiceBusConnectionString = "...";
        const string QueueName = "whizlabqueue";
        static IQueueClient whizlabqueueClient;
        static async Task SideSync()
        {
            whizlabqueueClient = new    Slot1    (ServiceBusConnectionString,
QueueName);
            await SendMessagesAsync(numberOfMessages);
            await queueClient.CloseAsync();
        }

        static async Task SendMessagesAsync(string messageBody)
        {
            try
           {
                    var message = new
Message(Encoding.UTF8.GetBytes(messageBody));
                    await queueClient.    Slot2    (message);
                }
            }
            catch (Exception whizlabexception)
            {
                Console.WriteLine($"{DateTime.Now} :: Exception:
{whizlabexception.Message}");
            } } } }
```

Which of the following needs to go into Slot2?

A. Send

B. SendMessage

✓ C. SendAsync

D. GetAsync

Explanation:

Answer – C

To send a message asynchronously, we have to use the “SendAsync” method.

An example of this is given in the Microsoft documentation

5. Directly after the `MainAsync()` method, add the following `SendMessagesAsync()` method that performs the work of sending the number of messages specified by `numberOfMessagesToSend` (currently set to 10):

C# Copy

```
static async Task SendMessagesAsync(int numberOfMessagesToSend)
{
    try
    {
        for (var i = 0; i < numberOfMessagesToSend; i++)
        {
            // Create a new message to send to the queue.
            string messageBody = $"Message {i}";
            var message = new Message(Encoding.UTF8.GetBytes(messageBody));

            // Write the body of the message to the console.
            Console.WriteLine($"Sending message: {messageBody}");

            // Send the message to the queue.
            await queueClient.SendAsync(message);
        }
    }
    catch (Exception exception)
    {
        Console.WriteLine($"{DateTime.Now} :: Exception: {exception.Message}");
    }
}
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with queues in .Net, one can go to the below URL

https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-dotnet-get-started-with-queues

Ask our Experts Rate this Question?

View Queries open

Question 55 Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

You have an Azure key vault in place named whizlabvault9000. You have to add a secret to the vault. Which of the following commands would you issue for this requirement?

A. az secret value set

✓ B. az keyvault secret set

C. az secret set

D. az secret key set

Explanation:

Answer – B

You have to issue the "az keyvault secret set" command.

An example is also given in the Microsoft documentation

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with secrets in the key vault, please visit the following URL

https://docs.microsoft.com/en-us/azure/key-vault/quick-create-cli

Ask our Experts

Rate this Question?

View Queries open

Finish Review

Certification

Cloud Certification

Java Certification

PM Certification

Big Data Certification

Company

Become Our Instructor

Support

Discussions

Blog

Business

Follow us

© Copyright 2020. Whizlabs Software Pvt. Ltd. All Right Reserved.